

## lecture 20

### Image Compositing

- chroma keying
- alpha
- F over B
- OpenGL blending
- chroma keying revisited: "pulling a matte"

### Organization of Course

- 1: Viewing transformations
- 2: Visibility, geometry modelling
- 3: Rendering: light, material, texture, **transparency**  
Transparency is a mix of rendering and image capture/display. It is a bridge between parts 3 and 4 of the course.
- 4: Image Capture and Display

Many computer graphics techniques use real images in some way.


- We have seen several examples
- scanned 3D models
  - texture mapping using photos
  - environment mapping

Let's start today's lecture with another example.

### Image Segmentation

Classic computer (and human) vision problem:

Partition an image into regions. It is a difficult problem (and not so well defined).



<http://www.eecs.berkeley.edu/Research/Projects/CS/vision/group1/resources.html>

Specific version of segmentation:

Given an image, partition it into a foreground and a background.



input



foreground

[http://www.cc.gatech.edu/~dellaert/07F-Vision/Schedule\\_files/10-LazySnapping.ppt.pdf](http://www.cc.gatech.edu/~dellaert/07F-Vision/Schedule_files/10-LazySnapping.ppt.pdf)

Computer graphics application: the foreground can then be pasted over a different background ("compositing")



input



(semi) automatic segmentation



output (composite with new background)

This is an old idea e.g. chroma-keying  
(green or blue screen)



It doesn't always work. (see video link)

<http://www.10tv.com/content/stories/2014/03/17/fray-fornsend-wears-green-disappears.html>



### General Approach

Step 1: Take picture of background B (not necessarily green screen)

Step 2: Take image/video of foreground character in front of background (F over B)

Step 3: // Compute foreground mask

For each pixel,  
if (F over B)(x,y) == B(x,y)  
mask(x,y) = 0 // background  
else mask(x,y) = 1 // foreground

Step 4: // Write foreground image over a new background Bnew

For each pixel (x,y)  
if mask(x,y) == 1  
I(x,y) = F(x,y)  
else I(x,y) = Bnew(x,y)

### Why doesn't it always work?

- Shadows (foreground object can change background)
- Interreflections (green screen can reflect, so foreground takes on color of background)
- Foreground object might happen to have same color as background (in step 3)
- Soft edges become hard (mask) e.g. Hair and furry object boundaries are difficult to model with a binary mask.

Ok, now let's look at a more general situation.

## lecture 20

### Image Compositing

- chroma keying
- alpha
- F over B
- OpenGL blending
- Chroma keying revisited: "pulling a matte"

### Partially occupied pixels & "alpha"

Think of a pixel as a little square. The occupancy or coverage of a pixel is called "alpha".

$\alpha = 0$  means not occupied at all (transparent).

$\alpha = 1$  means fully occupied (opaque)

$0 < \alpha < 1$  means partially occupied

In representing RGB images is common to include a 4th component to indicate how much of the pixel is occupied, so we have RGBA. Typically one uses 8 bits for each "channel" so this gives 32 bits per pixel.

## Examples of RGBA

(0, 0, 0, 1) - black and opaque  
 (1, 0, 0, 1) - red and opaque  
 etc.  
 (1, 1, 1, 1) - white and opaque  
 (.5, 0, 0, .5) - red and 50% transparent  
 (.5, .5, .5, .5) - white and 50% transparent  
 (.1, .1, .1, .5) - dark grey and 50% transparent  
 (.1, .1, .1, .1) - white and 10% transparent  
 (0, 0, 0, 0) - color undefined, 100% transparent

## Where do alpha values come from ?

In OpenGL, we can define surfaces as partially transparent.

e.g.

```
diffuse_material = [ 1, 0, 0, 0.5 ]
glMaterial(GL_FRONT, GL_DIFFUSE, diffuse_material)
drawPolygon()
```

The material has a red color with **50% transparency**.

I will sometimes write RGB and sometimes rgb.  
 The reasons will be explained later ("premultiplied values")

To give you a flavour of what's to come....

Q: How do we darken a pixel without changing its opacity ?


A:  $\text{darken}(I_{\text{rgba}}, \phi) = (\phi I_r, \phi I_g, \phi I_b, I_\alpha)$

Q: How do we change the opacity ( $\alpha$ ) of a pixel without changing the underlying color ?

$\text{dissolve}(I_{\text{rgba}}, \delta) = (\delta I_r, \delta I_g, \delta I_b, \delta I_\alpha)$

In the previous example, all triangles were in the  $z=0$  plane (and depth buffering was turned off). I just wanted to illustrate that the drawing order matters.

Here is another example which illustrates a more subtle point. For this example, there is no correct order to draw the two rectangles, since you cannot say that one rectangle is over another.



<http://stackoverflow.com/questions/1677432/opengl-alpha-blending-and-subject-independent-transparency>

If you draw blue first, then green will be drawn over blue at each pixel. However, there are some pixels in which the green rectangle is behind the blue one. Drawing the green first creates a similar problem, obviously.  
 The solution is similar to the painter's algorithm: split one of the rectangles and draw them from far to near.

## lecture 20

### Image Compositing

- chroma keying
- alpha
- F over B
- OpenGL blending
- Chroma keying revisited: "pulling a matte"

```
// glEnable(GL_BLEND)
// glBlendFunc(GL_SRC_ALPHA, GL_ONE_MINUS_SRC_ALPHA) // explain later
// glDisable(GL_DEPTH_TEST)
```

```
def drawYellowTriangle():
    glBegin(GL_TRIANGLES)
    glColor4f(1.0, 1.0, 0.0, 0.75) # yellow
    glVertex3f(0.1, 0.9, 0.0)
    glVertex3f(0.1, 0.1, 0.0)
    glVertex3f(0.7, 0.5, 0.0)
    glEnd()
```

```
def drawCyanTriangle():
    glBegin(GL_TRIANGLES)
    glColor4f(0.0, 1.0, 1.0, 0.75) # cyan
    glVertex3f(0.9, 0.9, 0.0)
    glVertex3f(0.9, 0.1, 0.0)
    glVertex3f(0.3, 0.5, 0.0)
    glEnd()
```

```
def drawMain():
    glPushMatrix()
    drawYellowTriangle() // right pair
    drawCyanTriangle()
```

```
glTranslatef(-1.0, 0)
drawCyanTriangle()
drawYellowTriangle() // left pair
glPopMatrix()
```

<https://www.opengl.org/archives/resources/code/samples/redbook/chapter18.html>



## F over B

Let's look at the "over" operation more formally.

How to put a foreground RGBA layer over a background RGBA layer?

I will use lower case "rgb" instead of RGB (for reasons to be explained later).

Notation: Foreground  $F_{rgb\alpha}$   
Background  $B_{rgb\alpha}$

Goal: How to compute a new RGBA layer which is the foreground layer over the background layer, i.e.

$$(F \text{ over } B)_{rgb\alpha} = ?$$

I changed the slide order and content a bit, relative to lecture.

### More general case:


background may be partly transparent,  $0 \leq B\alpha \leq 1$   
foreground may be partly transparent,  $0 \leq F\alpha \leq 1$

Again, given  $F_{rgb\alpha}$ ,  $B_{rgb\alpha}$ , how do we define  $(F \text{ over } B)_{rgb\alpha}$  ?

Note this is a per-pixel definition.

### Special but common case (opaque background):

background is opaque,  $B\alpha = 1$   
foreground may be partly transparent,  $0 \leq F\alpha \leq 1$



$$(F \text{ over } B)\alpha = F\alpha + (1 - F\alpha) B\alpha = 1$$

Let's not write out color yet.


### Example

Suppose the background color is black. Its RGB color is (0, 0, 0).

Suppose the foreground surface color is red.

We think of its RGB color as (1,0,0), e.g. `glColor(1, 0, 0)`

Suppose the surface has  $\alpha = 0.5$ .



There are two ways to view a partially occupied pixel. First, the pixel is transparent. Second, the underlying surface may be opaque but it only covers part of the pixel because it is near the boundary of the surface. For the present discussion, we don't care which of these two situations is present.

How should the RGBA values of the foreground pixel be interpreted/defined/represented ?

- You might argue it should be represented as

(1, 0, 0, 0.5)

since we have a red surface and the alpha value is 0.5.

- Or, you might argue that it should be represented as

(0.5, 0, 0, 0.5)

since the RGB value to be displayed at that pixel is (0.5, 0, 0).

Both are possible.

### Pre-multiplied color

In the latter case, (0.5, 0, 0, 0.5), we say the rgb values have "pre-multiplied" by  $\alpha$

$$(r, g, b, \alpha) = (\alpha R, \alpha G, \alpha B, \alpha)$$


RGB is the color that is computed when rendering e.g. with Blinn-Phong or `glColor()`. The value  $\alpha$  is given in the definition of the surface material or in `glColor()` as in our earlier example with cyan and yellow triangles.

[ASIDE: Note the similarity to homogeneous coordinates. e.g. (w x, w y, w z, w) represents the 3D point (x, y, z). ]

Given  $F_{rgb\alpha}$ ,  $B_{rgb\alpha}$ ,

how do we define  $(F \text{ over } B)_{rgb\alpha}$  ?

Assume the geometry below within a pixel. This would give us the formula below for the alpha values of the resulting layer.




$$(F \text{ over } B)_\alpha = F_\alpha + (1 - F_\alpha) B_\alpha$$

If we use pre-multiplied color values,

then we get the same formula:

[TODO: explain in lecture notes and/or exercises.]

$$(F \text{ over } B)_{rgb} = F_{rgb} + (1 - F_\alpha) B_{rgb}$$



You can think of this as:

$$(F \text{ over } B)_{rgb} = F_\alpha F_{RGB} + (1 - F_\alpha) B_\alpha B_{RGB}$$

Exercise: if we don't use premultiplied values, then we get a more complicated formula:

$$(F \text{ over } B)_{RGB} = \frac{F_\alpha F_{RGB} + (1 - F_\alpha) B_\alpha B_{RGB}}{F_\alpha + (1 - F_\alpha) B_\alpha}$$

We are distinguishing two representations:

- RGBA surface/object properties that you declare in OpenGL / etc
- Here, material & color are *independent* (which is preferable from the programmer's perspective).
- pre-multiplied pixel color values, rgba, that are written in the image buffer


We are distinguishing two representations:

- RGBA surface/object properties that you declare in OpenGL / etc
- Here, material & color are *independent* (which is preferable from the programmer's perspective).
- In terms of the graphics pipeline, this is a vertex property.
- pre-multiplied pixel color values, rgba, that are written in the image buffer

The transformation between the two happens in the fragment shader.

## OpenGL Blending


Blending must be enabled, else alpha is ignored and incoming fragment is written over the current pixel.




The fragment processor takes in fragments and uses them to modify pixels in the frame buffer i.e. image.



The fragment processor takes a fragment, and "blends" it with the current pixel to produce a modified pixel.




## Example 2: Blending not enabled (default)



This gives the new background.

## Example 1: (F over B)<sub>rgbα</sub>



This gives the new background.

Classic OpenGL offers several blending functions.

<http://www.khronos.org/opengl/docs/appendix6.html#name1>

glBlendFunc( source\_blending\_factor, destination\_blending\_factor )

For Example 1:

glBlendFunc(GL\_SRC\_ALPHA, GL\_ONE\_MINUS\_SRC\_ALPHA).

For Example 2:

glBlendFunc(GL\_ONE, GL\_ZERO).

Modern OpenGL allows you to write your own blending functions.

## Blending of Image Layers in Adobe Photoshop

This URL was recommended by the "orange book" OpenGL Shading Language.

See also [http://en.wikipedia.org/wiki/Blend\\_modes](http://en.wikipedia.org/wiki/Blend_modes)



B F

## lecture 20

### Image Compositing

- Chroma keying
- alpha
- F over B
- OpenGL blending
- Chroma keying revisited: "pulling a matte"

## "Pulling a matte"

(alpha channel = a "matte", binary alpha channel = a "mask")

We are given  $(F \text{ over } B)_{\text{rgba}}$  and maybe something else.

We would like to

- compute  $F_{\text{rgba}}$
- given a new new background  $B'$ ,  
compute  $(F \text{ over } B')_{\text{rgba}}$

## Alpha estimation using computer vision

Use one image only !

Exercise: Show you have 7 unknown variables at each pixel (but only 3 knowns, namely RGB).

Method:

Assume: F and B have non-overlapping different distributions of colors in 3D color space.

Allowed: user marks by hand regions that are B and other regions that are in F (and regions that may be in either).

This partitions the image pixels into three regions, called a **"tri-map"**.

[Ruzon and Tomasi, 2000]

original



"tri-map"



composed result on  
new background  
(matte not given)



Figure 11. A waterfall is transported to arid Death Valley.


## A Related Application: "1st and Ten"

<http://www.sportvision.com/>

<http://www.sportvision.com/media/1st-and-ten%E2%84%A2-line-system>



Exercise: what must be computed for this to work?



[Wang and Cohen 2006]