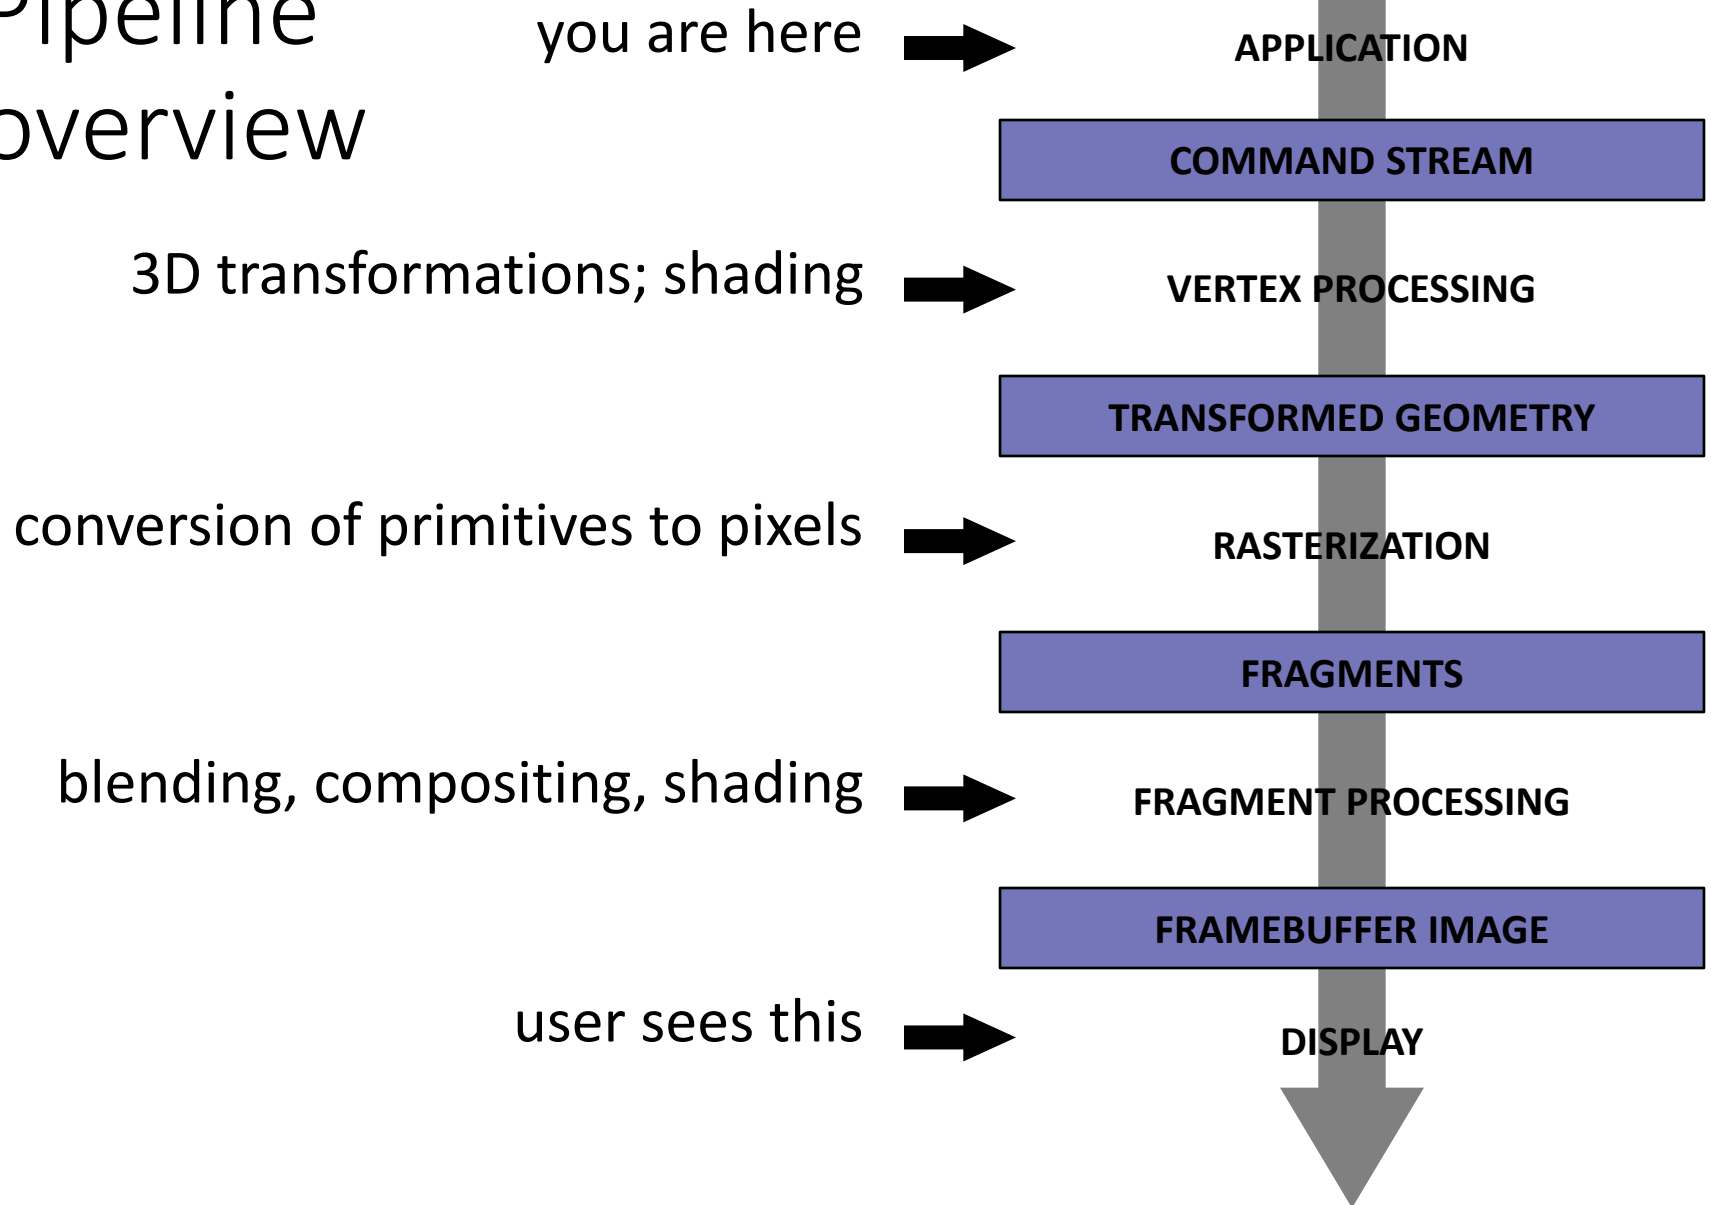


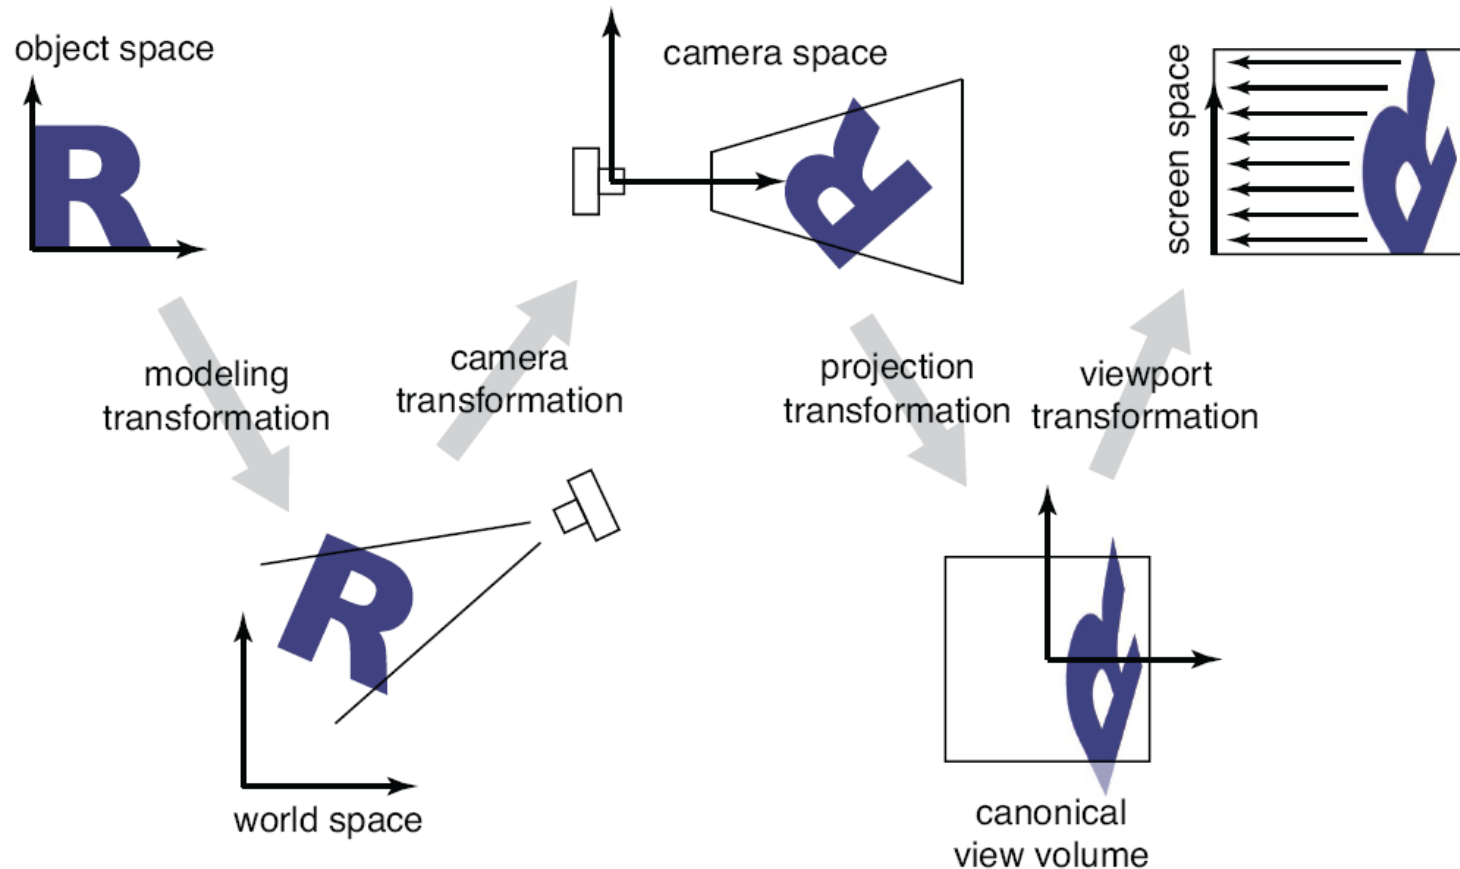
Pipeline Operations

Pipeline overview




Pipeline of transformations

- Standard sequence of transforms




Hidden surface elimination

- We have discussed how to map primitives to image space
 - projection and perspective are depth cues
 - occlusion is another very important cue



Back face culling


- For closed shapes you will never see the inside
 - therefore only draw surfaces that face the camera
 - ~~**Could** implement by checking $\mathbf{n} \cdot \mathbf{v}$~~
~~but \mathbf{v} varies across the surface (might want \mathbf{n} to vary too!)~~
 - **Actually** implemented by checking counter clockwise order for front facing triangles in *screen space*



USE A 3D drawing


To show front faces and back faces

This $\mathbf{n} \cdot \mathbf{v}$ explanation is not helpful




Painter's algorithm

- Simplest way to do hidden surfaces
- Draw from back to front, use overwriting in framebuffer




Painter's algorithm

- Useful when a valid order is easy to come by
- Compatible with alpha blending



[Foley et al.]


Drawing




Projection (left) and rasterization (right) of a triangle.

Visibility

- Painter's algorithm
 - Sort back to front
 - Draw!




- Doesn't always work... ☹️




Visibility

- Warnock's algorithm
 - Area subdivision
 - Apply Painter's when it will work (e.g., individual pixels)




Visibility

- Binary space partition
 - Linear time back to front sort
 - Key to 3D games before consumer level GPUs (Doom 1993)



Visibility


- Z-buffer
 - Store depth at every pixel
 - Compare when rasterizing



The z buffer

- In many (most) applications maintaining a z sort is too expensive
 - changes all the time as the view changes
 - many data structures exist, but complex
- Solution: draw in any order, keep track of closest
 - allocate extra channel per pixel to keep track of closest depth so far
 - when drawing, compare object's depth to current closest depth and discard if greater
 - this works just like any other compositing operation

The z buffer




- another example of a memory-intensive brute force approach that works and has become the standard

Precision in z buffer

- The precision is distributed between the near and far clipping planes
 - this is why these planes have to exist
 - also why you can't always just set them to very small and very large distances
- Generally use z' (not world z) in z buffer

Interpolating in projection




linear interp. in screen space \neq linear interp. in world (eye) space

More precision close to near plane (can revisit projection demo)

Pipeline for minimal operation

- Vertex stage (input: position / vtx; color / tri)
 - transform position (object to screen space)
 - pass through color
- Rasterizer
 - pass through color
- Fragment stage (output: color)
 - write to color planes


Result of minimal pipeline



Pipeline for basic z buffer


- Vertex stage (input: position / vtx; color / tri)
 - transform position (object to screen space)
 - pass through color
- Rasterizer
 - interpolated parameter: z' (screen z)
 - pass through color
- Fragment stage (output: color, z')
 - write to color planes only if interpolated $z' < \text{current } z'$

Result of z-buffer pipeline



Flat shading

- Shade using the real normal of the triangle
 - same result as ray tracing a bunch of triangles
- Leads to constant shading and faceted appearance
 - truest view of the mesh geometry




[Foley et al.]

Pipeline for flat shading


- Vertex stage (input: position / vtx; color and normal / tri)
 - transform position and normal (object to eye space)
 - compute shaded color per triangle using normal
 - transform position (eye to screen space)
- Rasterizer
 - interpolated parameters: z' (screen z)
 - pass through color
- Fragment stage (output: color, z')
 - write to color planes only if interpolated $z' < \text{current } z'$

Result of flat-shading pipeline



Local vs. infinite viewer, light

- Phong illumination requires geometric information:
 - light vector (function of position)
 - eye vector (function of position)
 - surface normal (from application)
- Light and eye vectors change
 - need to be computed (and normalized) for each face




Local vs. infinite viewer, light

- Look at case when eye or light is far away:
 - distant light source: nearly parallel illumination
 - distant eye point: nearly orthographic projection
 - in both cases, eye or light vector changes very little
- Optimization: approximate eye and/or light as infinitely far away


Directional light

- Directional (infinitely distant) light source
 - light vector always points in the same direction
 - often specified by position $[x \ y \ z \ 0]$
 - many pipelines are faster if you use directional lights



Infinite viewer


- Orthographic camera
 - projection direction is constant
- “Infinite viewer”
 - even with perspective, can approximate eye vector using the image plane normal
 - can produce weirdness for wide-angle views
 - Blinn-Phong: light, eye, half vectors all constant!



Pipeline for Gouraud shading

- Vertex stage (input: position, color, and normal / vtx)
 - transform position and normal (object to eye space)
 - compute shaded color per vertex
 - transform position (eye to screen space)
- Rasterizer
 - interpolated parameters: z' (screen z); r, g, b color
- Fragment stage (output: color, z')
 - write to color planes only if interpolated $z' < \text{current } z'$


Result of Gouraud shading pipeline



Vertex normals

- Need normals at vertices to compute Gouraud shading
- Best to get vtx. normals from the underlying geometry
 - e. g. spheres example
- Otherwise have to infer vtx. normals from triangles
 - simple scheme: average surrounding face normals

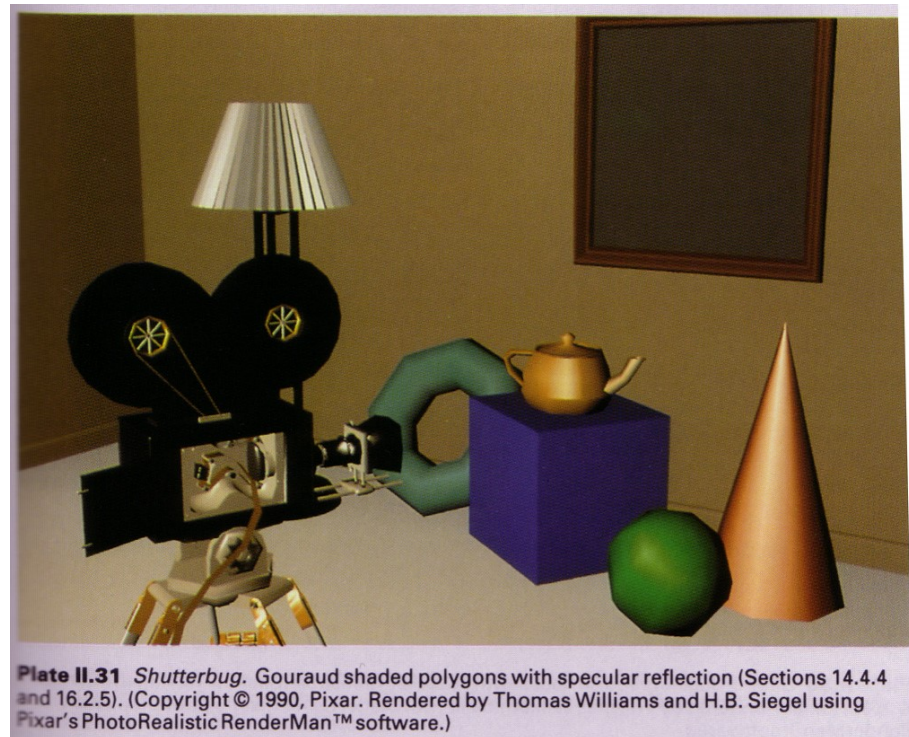
$$N_v = \frac{\sum_i N_i}{\|\sum_i N_i\|}$$



[Foley et al.]

Non-diffuse Gouraud shading


- Can apply Gouraud shading to any illumination model
 - it's just an interpolation method
- Results are not so good with fast-varying models like specular ones
 - problems with any highlights smaller than a triangle



[Foley et al.]

Phong shading

- Get higher quality by interpolating the normal
 - just as easy as interpolating the color
 - but now we are evaluating the illumination model per pixel rather than per vertex (and normalizing the normal first)
 - in pipeline, this means we are moving illumination from the vertex processing stage to the fragment processing stage



[Foley et al.]

Phong shading

- Bottom line: produces much better highlights



Shutterbug. Gouraud shaded polygons with specular reflection (Sections 14.4.4 and 16.2.5). (Copyright © 1990, Pixar. Rendered by Thomas Williams and H.B. Siegel using Pixar's PhotoRealistic RenderMan™ software.)

Plate II.32 Shutterbug. Phong shaded polygons with specular reflection (Sections 14.4.4 and 16.2.5). (Copyright © 1990, Pixar. Rendered by Thomas Williams and H.B. Siegel using Pixar's PhotoRealistic RenderMan™ software.)



[Foley et al.]

Pipeline for Phong shading

- Vertex stage (input: position, color, and normal / vtx)
 - transform position and normal (object to eye space)
 - transform position (eye to screen space)
 - pass through color
- Rasterizer
 - interpolated parameters: z' (screen z);
 r, g, b color; x, y, z **normal**
- Fragment stage (output: color, z')
 - compute shading using interpolated color and normal
 - write to color planes only if interpolated $z' < \text{current } z'$

Result of Phong shading pipeline

