

Ray Tracing II

SOCS Teaching Assistant Excellence Award

- Had a great experience with one of your TAs?
- Nominate them for a TA Award!

Nomination Information:

<https://www.cs.mcgill.ca/academic/ta/taawards>

Criteria: organized, knowledgeable, effective, accessible, inspiring, role model...





Topics

- Transformations in ray tracing
 - Transforming objects
 - Transformation hierarchies
- Ray tracing acceleration structures
 - Bounding volumes
 - Bounding volume hierarchies
 - Uniform spatial subdivision
 - Adaptive spatial subdivision

Transforming objects

- In modeling, we've seen the usefulness of transformations
 - How to do the same in RT?
- Take spheres as an example: want to support transformed spheres
 - Need a new Surface subclass
- Option 1: transform sphere into world coordinates
 - Write code to intersect arbitrary ellipsoids
- Option 2: transform ray into sphere's coordinates
 - Then just use existing sphere intersection routine

Intersecting transformed objects



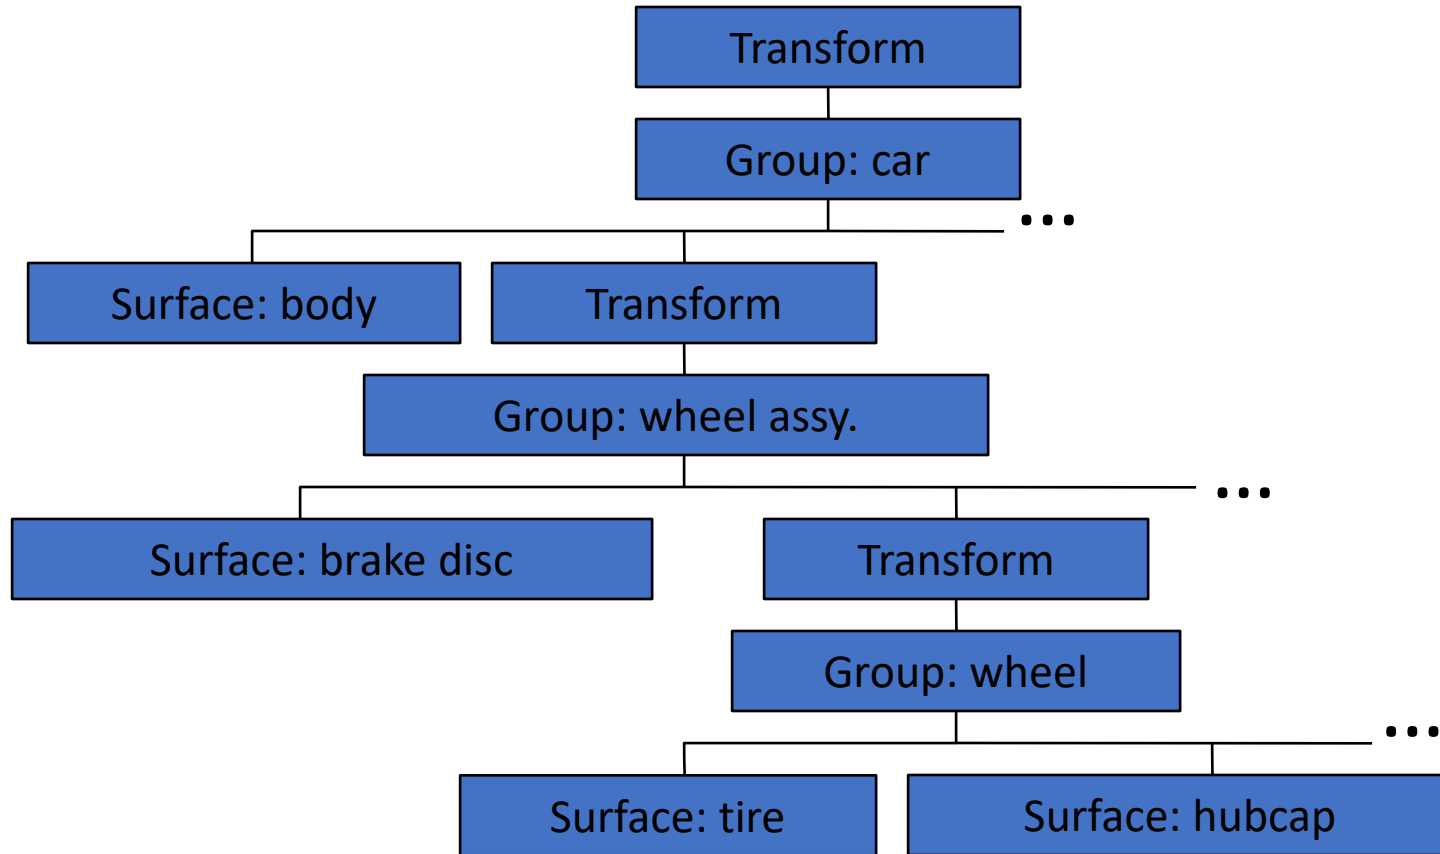
Implementing RT transforms

- Create wrapper object “TrasformedSurface”
 - Has a transform T and a reference to a surface S
- To intersect with a transformed surface:
 - Transform ray to local coords (by inverse of T)
 - Call `surface.intersect`
 - Transform hit data back to global coords (by T)
 - Intersection point
 - Surface normal
 - Any other relevant data (maybe none)

Groups, transforms, hierarchies

- Often it's useful to transform several objects at once
 - Add “SurfaceGroup” as a subclass of Surface
 - Has a list of surfaces
 - Returns closest intersection
 - Opportunity to move ray intersection code here to avoid duplication
- With TransformedSurface and SurfaceGroup you can put transforms below transforms
 - This gives us a transformation hierarchy.

A transformation hierarchy

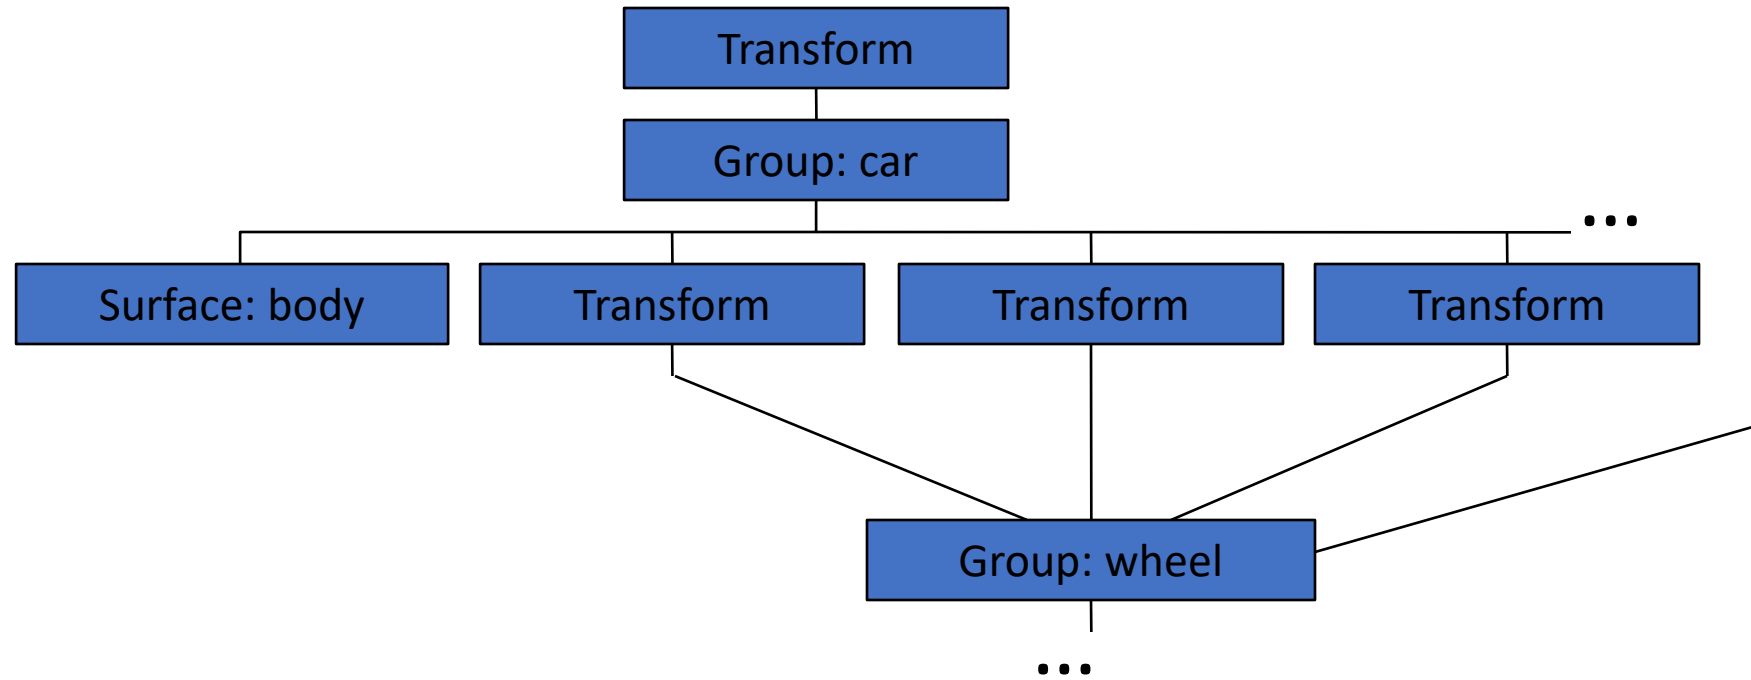


- Common optimization: merge transforms with groups
 - This is how we did it in the modeler assignment

Instancing

- Anything worth doing is worth doing n times
- If we can transform objects, why not transform them several ways?
 - Many models have repeated subassemblies
 - Mechanical parts (wheels of car)
 - Multiple objects (chairs in classroom, ...)
 - Nothing stops you from creating two TransformedSurface objects that reference the same Surface
 - Allowing this makes the transformation tree into a DAG
 - (directed acyclic graph)
 - Mostly this is transparent to the renderer

Hierarchy with instancing



Hierarchies and performance


- Transforming rays is expensive
 - minimize tree depth: flatten on input
 - push all transformations toward leaves
 - triangle meshes may do best to stay as group
 - transform ray once, intersect with mesh
 - internal group nodes still required for instancing
 - can't push two transforms down to same child!

Ray tracing acceleration

- Ray tracing is slow. This is bad!
 - Ray tracers spend most of their time in ray-surface intersection methods
- Ways to improve speed
 - Make intersection methods more efficient
 - Yes, good idea. But only gets you so far
 - Call intersection methods fewer times
 - Intersecting every ray with every object is wasteful
 - Basic strategy: efficiently find big chunks of geometry that definitely do not intersect a ray

Bounding volumes

- Quick way to avoid intersections: bound object with a simple volume
 - Object is fully contained in the volume
 - If it doesn't hit the volume, it doesn't hit the object
 - So test bvol first, then test object if it hits



[Glassner 89, Fig 4.5]

Bounding volumes

- Cost:
 - slightly more for hits and near misses,
 - ***much less*** for far misses
- Worth doing? It depends:
 - Cost of bvol intersection test should be small
 - Therefore use simple shapes (spheres, boxes, ...)
 - Cost of object intersect test should be large
 - Bvols most useful for complex objects
 - Tightness of fit should be good
 - Loose fit leads to extra object intersections
 - Tradeoff between tightness and bvol intersection cost

Implementing bounding volume


- Just add new Surface subclass, “BoundedSurface”
 - Contains a bounding volume and a reference to a surface
 - Intersection method:
 - Intersect with bvol, return false for miss
 - Return surface.intersect(ray)
 - Like transformations, common to merge with group
 - This change is transparent to the renderer (only it might run faster)
- Note that all Surfaces will need to be able to supply bounding volumes for themselves

If it's worth doing, it's worth doing hierarchically!


- Bvols around objects may help
- Bvols around groups of objects will help
- Bvols around parts of complex objects will help
- Leads to the idea of using bounding volumes all the way from the whole scene down to groups of a few objects
- Meshes are a special case
 - We don't add individual triangles into the DAG
 - We probably want to do something hierarchically at the level of an individual triangle soup!

Implementing a bvol hierarchy


- A BoundedSurface can contain a list of Surfaces
- Some of those Surfaces might be more BoundedSurfaces
 - Axis Aligned Bounding Boxes (AABB), Easy!
 - Spheres, can computed in expected $O(n)$ time
 - WELZL, E. 1991. Smallest enclosing disks (balls and ellipsoids)
 - Other fast approximate methods
 - Oriented Bounding Boxes (OBBs), and other types also possible



BVH construction example



BVH ray-tracing example



Choice of bounding volumes

- Spheres
 - easy to intersect, not always so tight,
 - Tricky to compute tight bounds (google Welzl's smallest enclosing disk algorithm)
- Axis-aligned bounding boxes (AABBs) –
 - easy to intersect, often tighter (esp. for axis-aligned models)
- Oriented bounding boxes (OBBs) –
 - easy to intersect (but cost of transformation), tighter for arbitrary objects
- Computing the bvols
 - For primitives -- generally pretty easy
 - For groups -- not so easy for OBBs (to do well)
 - For transformed surfaces -- not so easy for spheres

Axis aligned bounding boxes

- Probably easiest to implement
- Computing for primitives
 - Cube: very straightforward!
 - Sphere, cylinder, etc.: easy to approximate
 - Groups or meshes: min/max of component parts
- AABBs for transformed surface
 - Easy to do conservatively: bbox of the 8 corners of the bbox of the untransformed surface
- Can compute intersection as we saw previously as the intersection of 3 slabs

Building a hierarchy


- Usually do it top-down
- Make bbox for whole scene, then split into (maybe 2) parts
 - Recurse on parts
 - Stop when there are just a few objects in your box

Building a hierarchy

- How to partition?
 - Ideal: clusters
 - Practical: partition along axis
 - Median partition
 - More expensive
 - More balanced tree
 - Center partition
 - Less expensive, simpler
 - Unbalanced tree, but that may actually be better


Regular space subdivision

- An entirely different approach: uniform grid of cells






Regular grid example

- Grid divides space, not objects




Traversing a regular grid



Non-regular space subdivision

- *k*-d Tree
 - subdivides space, like grid
 - adaptive, like BVH



Implementing acceleration structures


- Conceptually simple to build acceleration structure into scene structure
- Better engineering decision to separate them

Other ray tracing features

- Fresnel reflection and refraction
- Area light sources / soft shadows
- Super sampling for anti aliasing & jittering
- Depth of field blur / real lens modelling
- Motion blur
- Transparency and compositing
- Cartoon shaders
- Texture maps (colour, displacement, normal)
- Other implicit surfaces (metaballs)
- Subdivision surfaces / parametric surfaces
- Constructive solid geometry
- Perlin Noise

Constructive Solid Geometry

- Exactly like intersecting a cube using 3 slabs
- Assemble binary set operations in a tree (works nicely with a dag scene graph)
- Ray intersection class must maintain a list of all intersections instead of closest intersection.




Meta Balls

- Implicit function (other (better) choices exist!)

$$\sum_{i=0}^n f_i(\mathbf{x}) \leq \text{threshold}$$

$$f_i(\mathbf{x}) = \frac{1}{\|\mathbf{x} - \mathbf{x}_i\|^2}$$

- ***Consider the picture at right***
 - ***What will be the normal?***
 - ***How do you find the intersection?***
 - ***Where do the colours come from?***



Meta Balls - Normals



- Not knowing the equation for $f_i(\mathbf{x})$ can we approximate $\nabla_{\mathbf{x}} f_i$?


- Finite differences?

- If $f_i(\mathbf{x}) = \frac{1}{\|\mathbf{x} - \mathbf{x}_i\|^2}$, then what is $\nabla_{\mathbf{x}} f_i$?

- For $r > 0$, let $p(r) = \begin{cases} 6r^5 - 15r^4 + 10r^3 & r < 1 \\ 0 & \text{otherwise} \end{cases}$



- Why is this nice? Finite support!

- If $f_i(\mathbf{x}) = p(\|\mathbf{x} - \mathbf{x}_i\|)$, then what is $\nabla_{\mathbf{x}} f_i$?

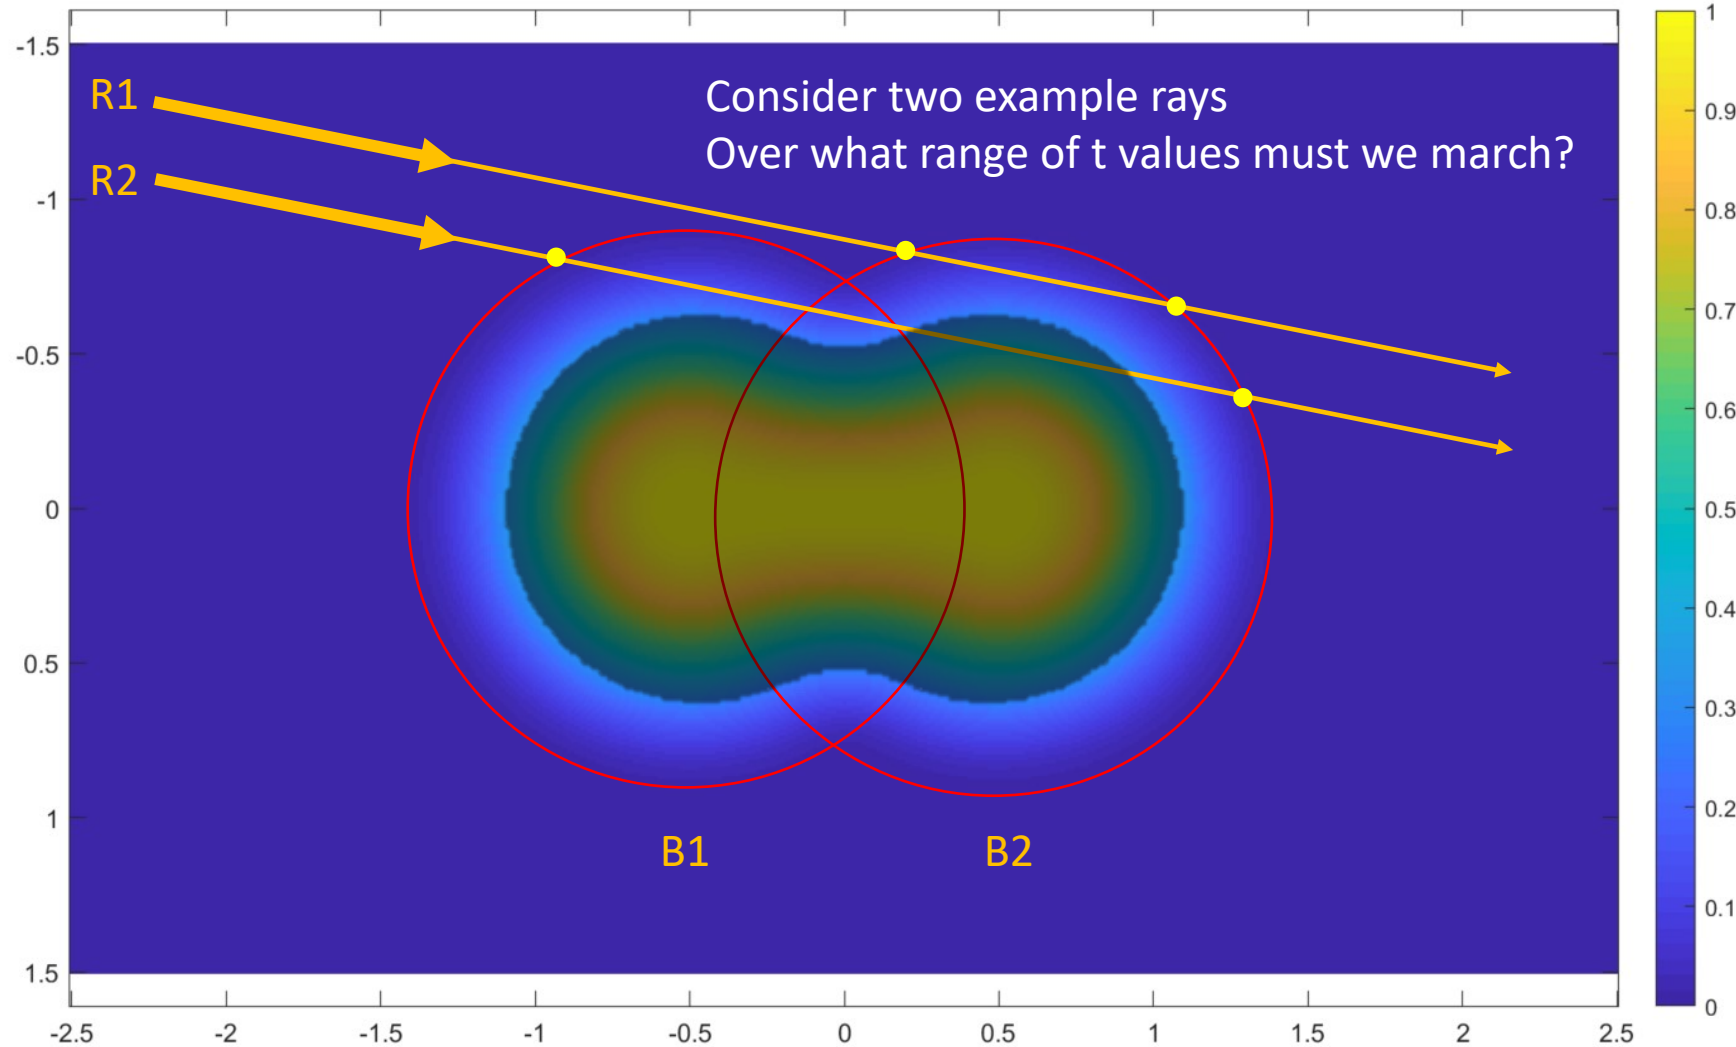


Meta Balls – Ray marching

- Easiest approach to meta ball intersection is ray marching
 - Given a signed distance function, then at any point we can compute the **distance** and know that we can step ***at least this far*** along the ray




Meta Balls – Ray marching



- Computing the exact distance? Tricky!
 - But can underestimate!
 - Use Lipschitz constant

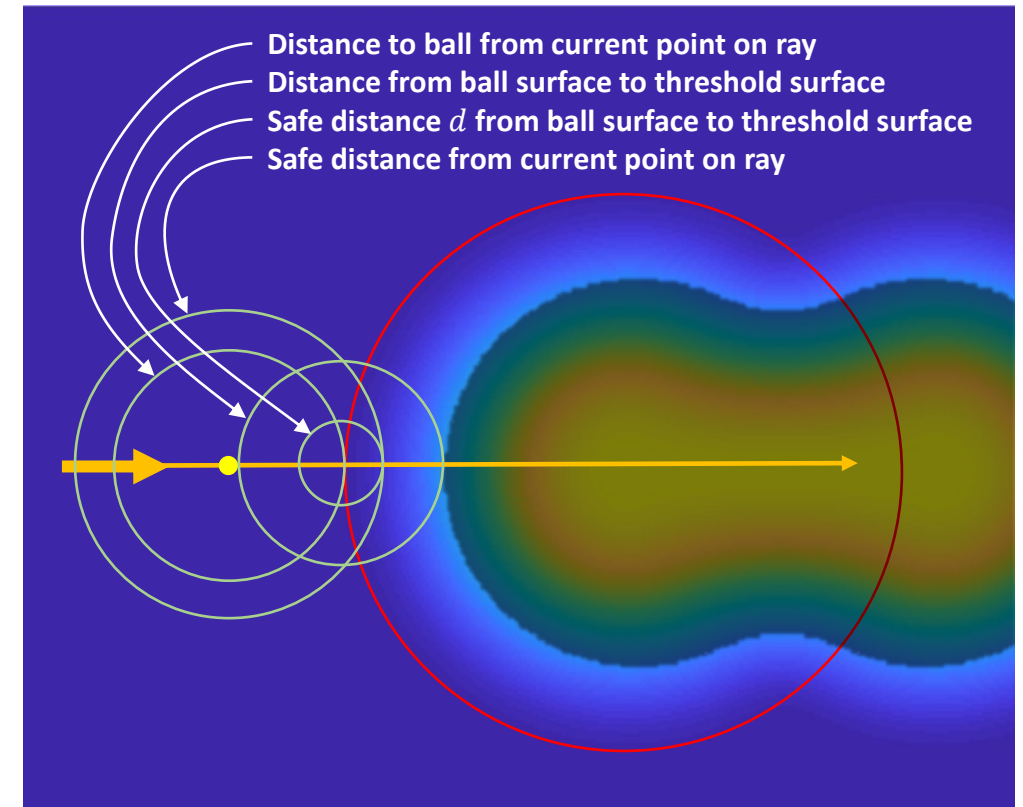

$$|p(r_1) - p(r_2)| \leq M|r_1 - r_2|$$

- Lipschitz constant M is the maximum of 1st derivative on $[0,1]$





Ray marching with Lipschitz bound

- The Lipschitz bound M gives us worst case growth, for which we can identify a safe underestimate of the distance d to the surface



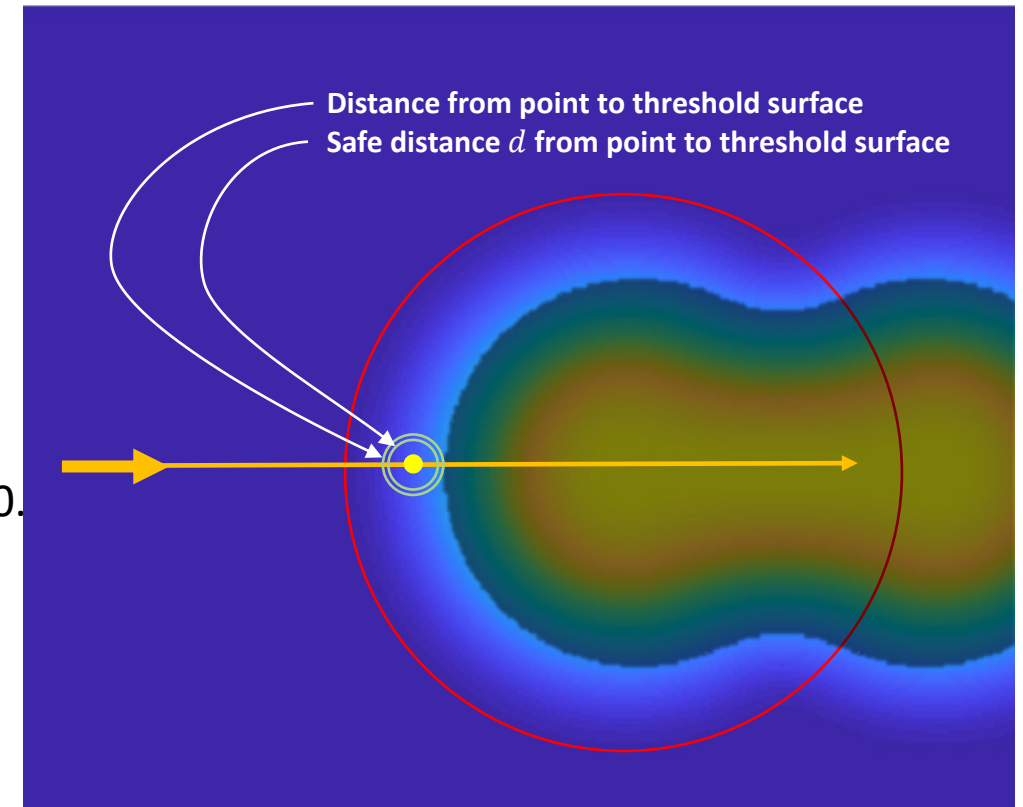

Ray marching with Lipschitz bound

- The Lipschitz bound M gives us worst case growth, for which we can identify a safe underestimate of the distance d to the surface




Ray marching with Lipschitz bound

- The Lipschitz bound M gives us worst case growth, for which we can identify a safe underestimate of the distance d to the surface



Ray marching – be careful!





- When the ray point falls into multiple balls, the implicit field function is the sum of both
- The Lipschitz bound is likewise the sum of the bounds.

Review and more information on sphere-tracing of implicit surfaces:

<https://www.scratchapixel.com/lessons/advanced-rendering/rendering-distance-fields>



Subdivision Surfaces

- Many subdivision schemes simply involve a **convex combination of points**
 - Chaikin, Lane-Reisenfeld, Loop, Catmull-Clark
- Convex hull of points can be used for ray-intersection test
 - Recursively subdivide checking with bounding box of all child faces
 - Once sufficiently flat (or small), use ray polygon intersection



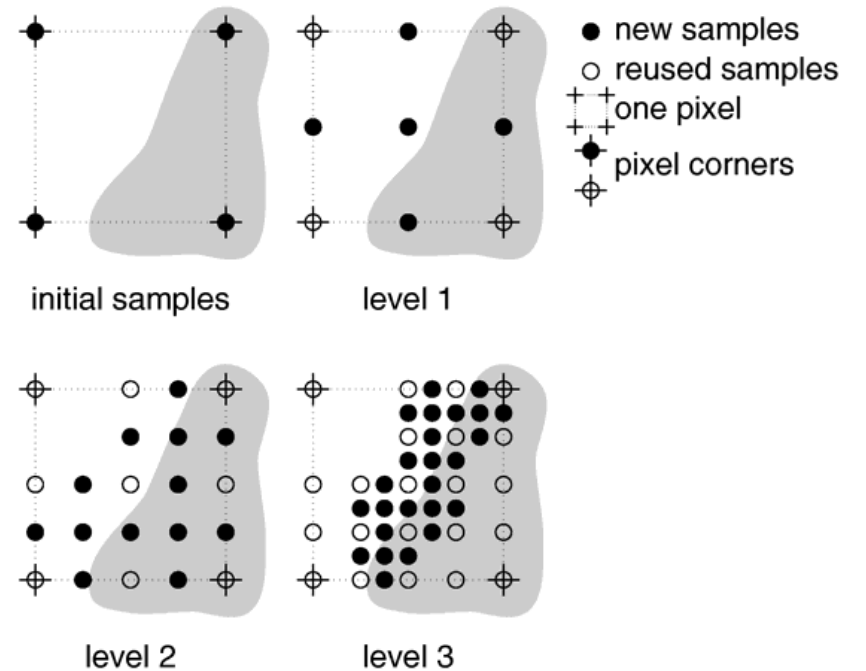

Perlin Noise

- Useful for clouds, marble, and other random varying effects
- Code is very small... noise is easy to compute!





Super Sampling and Jittering Samples

- Cast more rays through sub pixel positions.
- Watch out for aliasing!
 - Jittering samples will replace aliasing with noise, which is preferable
 - Noise added to sub-pixel position should not be too large!
- Super sampling can be done adaptively based on scene complexity at the pixel



Depth of Field, Area Lights

- Cast more rays!
 - Sample different ***random positions*** on lens aperture by changing the eye position
 - For an area light source, do shadow ray test and lighting computation for a different ***random positions*** on the area light source.



Motion Blur

- Knowing object trajectories over a time interval, cast rays at different times and average.
- Can also just average a sequence of images, but this would produce temporal aliasing!
- Can also be smart about not casting few rays where nothing is in motion.



[Thomas Porter 1984]



Reflection and Refraction



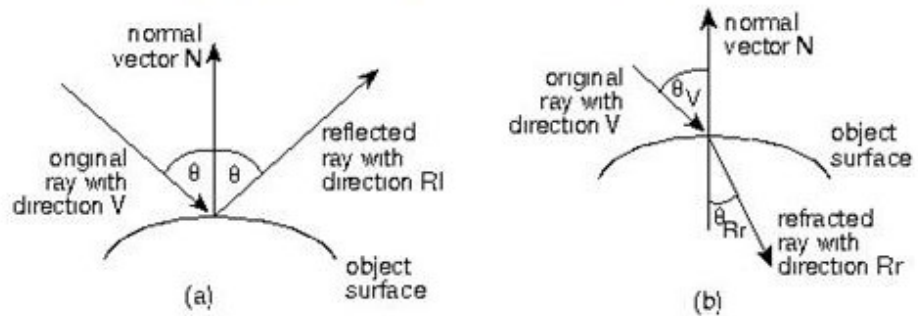
$n1$ = index of refraction of original medium

$n2$ = index of refraction of new medium

$$n = \frac{n1}{n2}$$


$$c2 = \sqrt{1 - n^2 \times (1 - c1^2)}$$

$$Rr = (n \times V) + (n \times c1 - c2) \times SN$$



Caustics

- Reflection and refraction scatters light too
 - Very challenging without photon mapping, or path tracing, etc.



Anisotropic BRDFs

- Can use a varying reflectance function for computing the amount of reflected light.




Cartoon Shaders

- Change lighting computations so that they produce a smaller number of colours.
- Add black edges at boundaries between objects.



Texture Maps

- Can download texture images from a variety of repositories.
- Easy to texture certain primitives (cubes, spheres, planes)
- Meshes often have texture coordinates attached
- ***Sampling is an issue!!***
 - Resolution may be too high or too low giving a poor quality result



Review and More Information

- FCG Chapters 12 and 13
- <https://www.scratchapixel.com/lessons/advanced-rendering/rendering-distance-fields>
- CGPP Chapter 27 Material scattering models
- CGPP Chapter 37 Spatial data structures