



PHYSICALLY-BASED RENDERING

Based on slides by Derek Nowrouzezahrai

Monte Carlo for Global Illumination

MC for the Rendering Equation

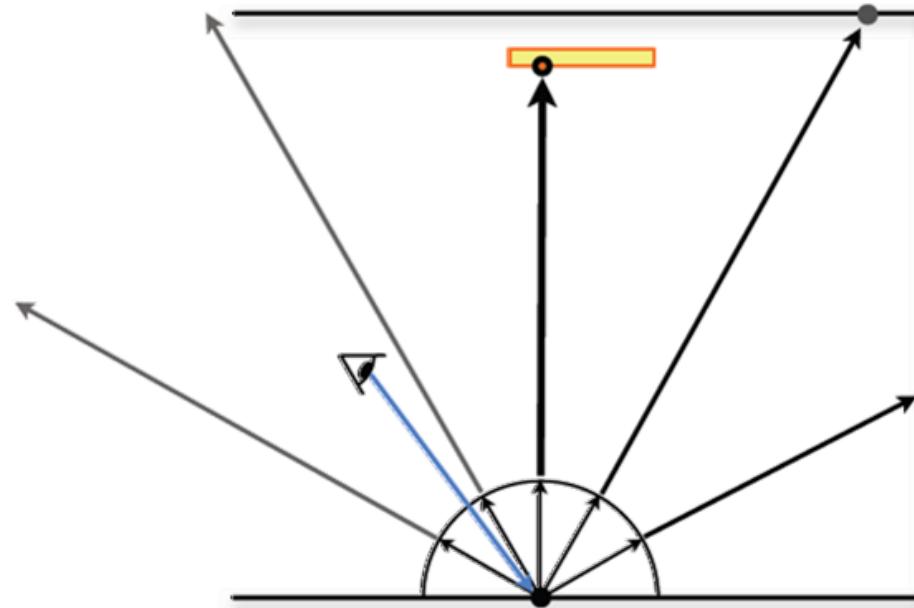
- What happens if we apply MC integration to the rendering equation?

$$L(x, \omega_o) = L_e(x, \omega_o) + \int_{\Omega} L(x', \omega_i) f_r(x, \omega_i, \omega_o) \cos \theta_i d\omega_i$$

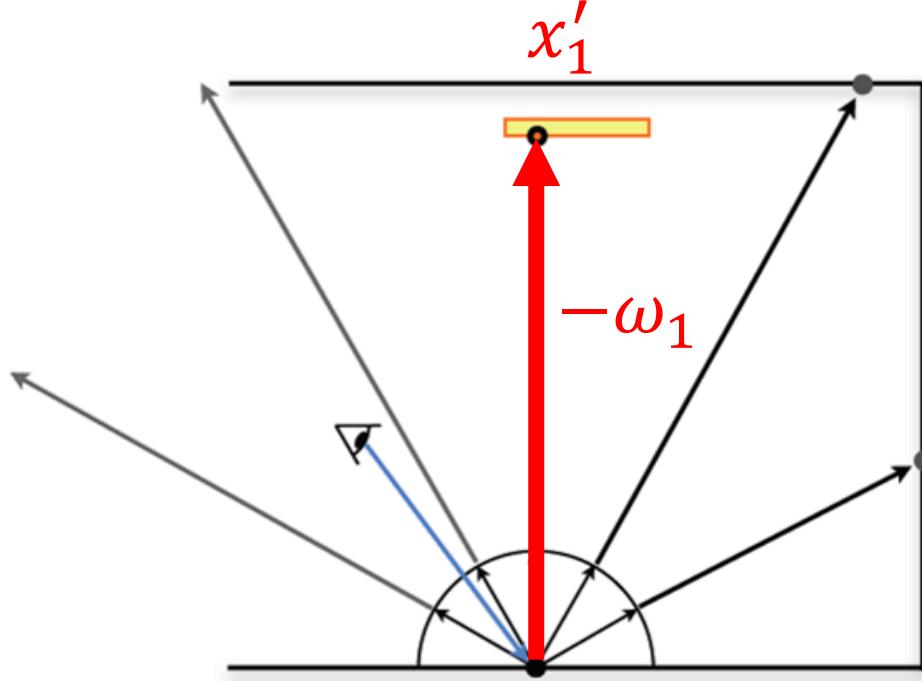
$$L(x, \omega_o) \approx L_e(x, \omega_o) + \frac{1}{N} \sum_i \frac{L(x'_i, \omega_i) f_r(x, \omega_i, \omega_o) \cos \theta_i}{pdf(\omega_i)}$$

MC for the Rendering Equation

$$L(x, \omega_o) \approx L_e(x, \omega_o) + \frac{1}{N} \sum_i \frac{L(x'_i, \omega_i) f_r(x, \omega_i, \omega_o) \cos \theta_i}{pdf(\omega_i)}$$



MC for the Rendering Equation



$$f_r(x'_1) = 0$$

$$L(x'_1, \omega_1) = L_e(x'_1, \omega_1)$$

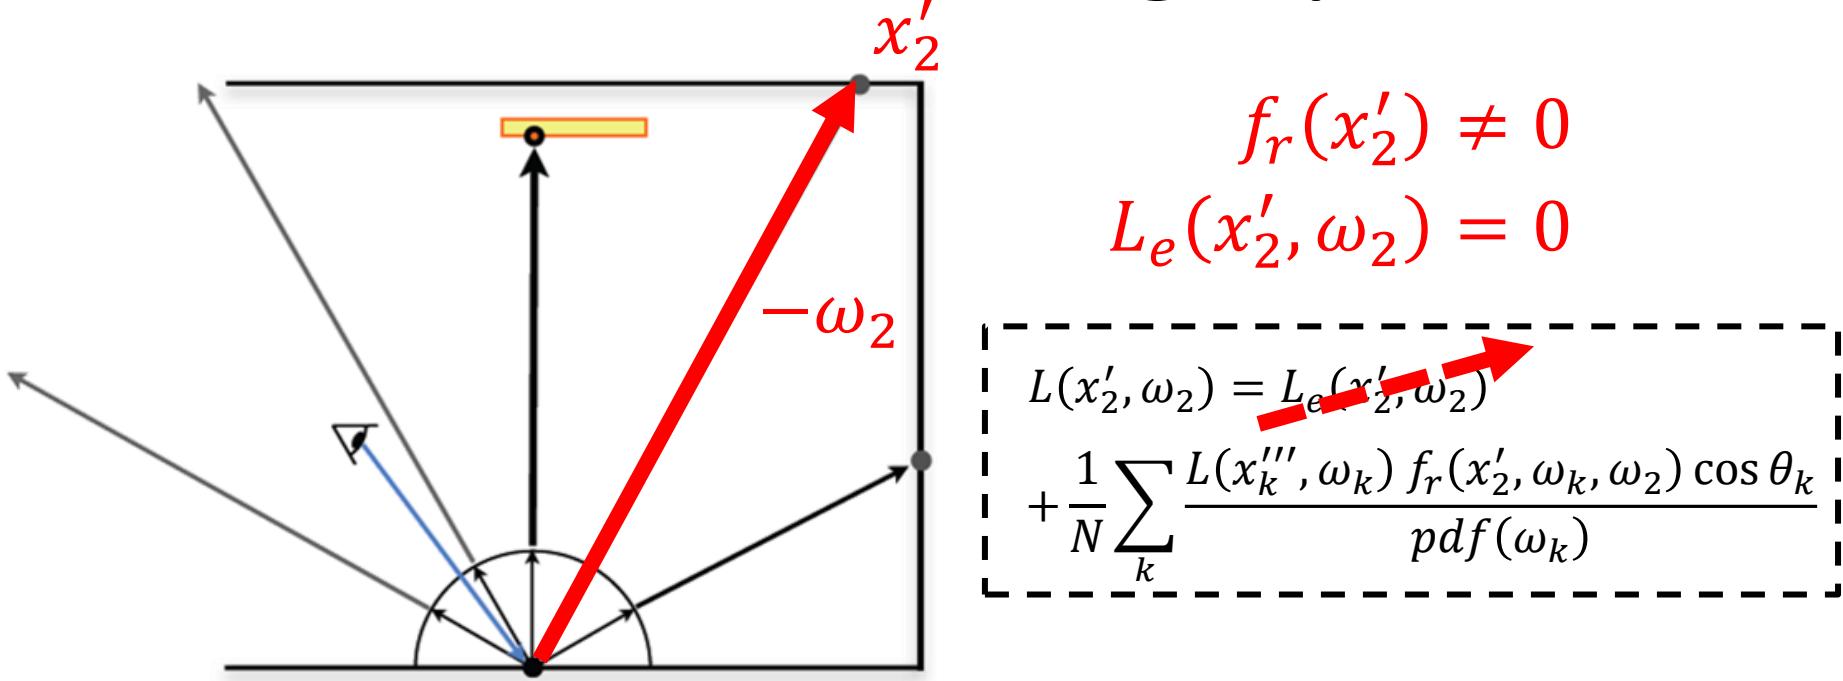
$$L(x'_1, \omega_1) = L_e(x'_1, \omega_1)$$

$$+ \frac{1}{N} \sum_j \frac{L(x''_j, \omega_j) f_r(x'_1, \omega_j, \omega_1) \cos \theta_j}{pdf(\omega_j)}$$

- What's the value at the sample $L(x'_1, \omega_1)$?

$$L(x, \omega_o) = L_e(x, \omega_o) + \frac{1}{N} \sum_i \frac{L(x'_i, \omega_i) f_r(x, \omega_i, \omega_o) \cos \theta_i}{pdf(\omega_i)}$$

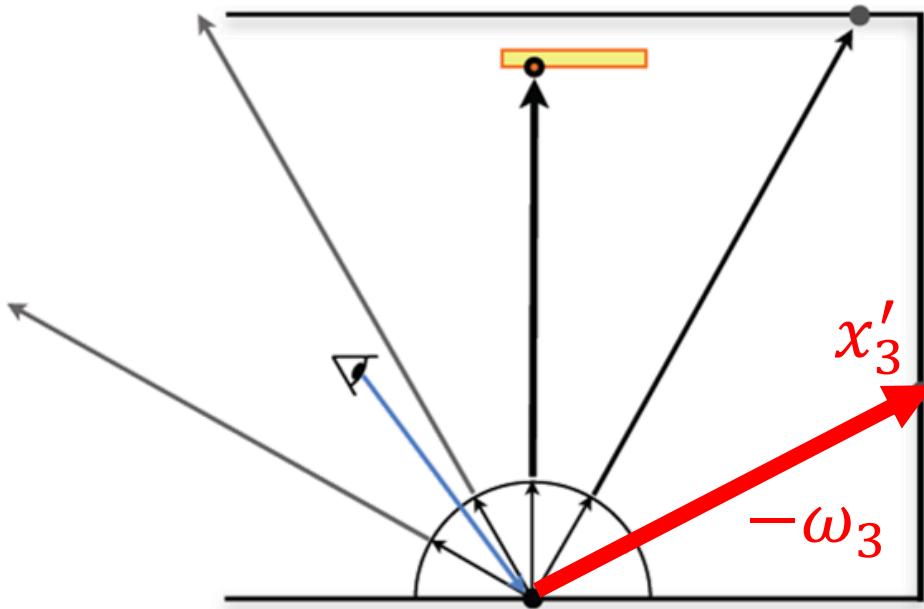
MC for the Rendering Equation



- What's the value at the sample $L(x'_2, \omega_2)$?

$$L(x, \omega_o) = L_e(x, \omega_o) + \frac{1}{N} \sum_i \frac{L(x'_i, \omega_i) f_r(x, \omega_i, \omega_o) \cos \theta_i}{pdf(\omega_i)}$$

MC for the Rendering Equation



$$f_r(x'_2) \neq 0$$

$$L_e(x'_2, \omega_2) = 0$$

$$f_r(x'_3) \neq 0$$

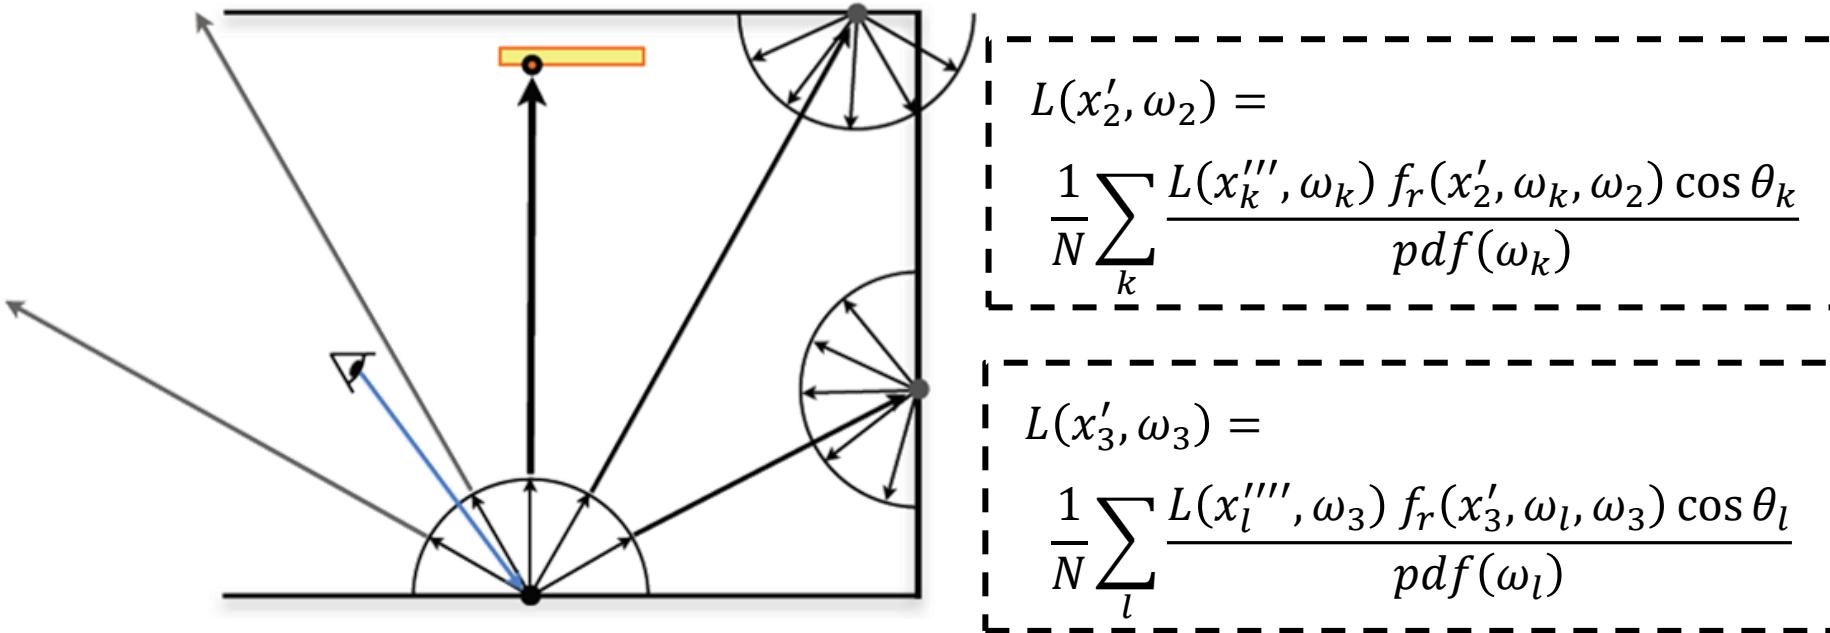
$$L_e(x'_3, \omega_3) = 0$$

$$\begin{aligned} L(x'_3, \omega_3) &= L_e(x'_3, \omega_3) \\ &+ \frac{1}{N} \sum_l \frac{L(x''', \omega_3) f_r(x'_3, \omega_l, \omega_3) \cos \theta_l}{pdf(\omega_l)} \end{aligned}$$

- What's the value at the sample $L(x'_3, \omega_3)$?

$$L(x, \omega_o) = L_e(x, \omega_o) + \frac{1}{N} \sum_i \frac{L(x'_i, \omega_i) f_r(x, \omega_i, \omega_o) \cos \theta_i}{pdf(\omega_i)}$$

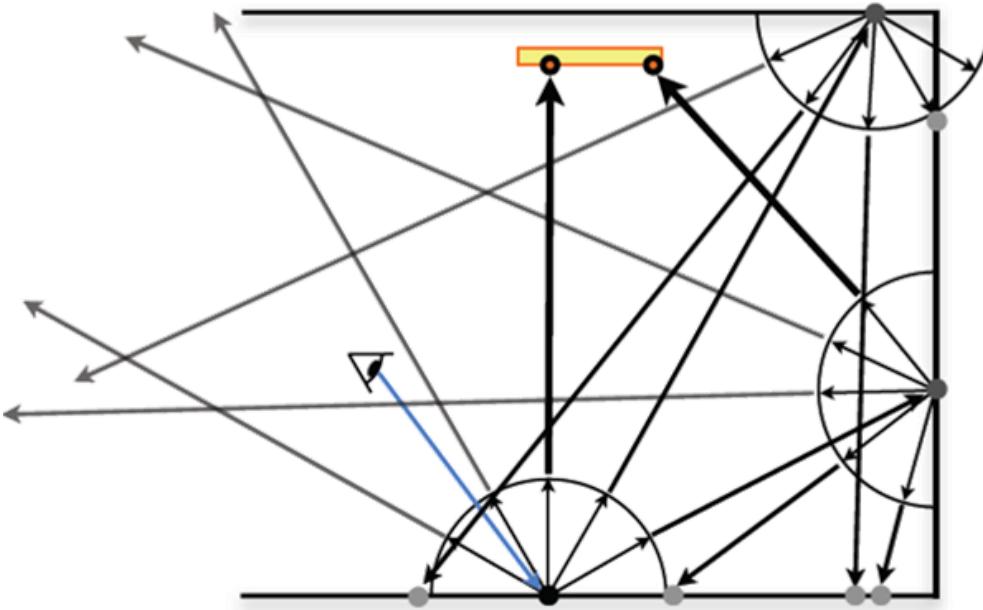
MC for the Rendering Equation



- And so on and so forth, proceeding recursively...

$$L(x, \omega_o) = L_e(x, \omega_o) + \frac{1}{N} \sum_i \frac{L(x'_i, \omega_i) f_r(x, \omega_i, \omega_o) \cos \theta_i}{pdf(\omega_i)}$$

MC for the Rendering Equation



```
color shade( point x )
  if( x.Le != 0 ) return x.Le;
  color L = 0;
  for( for sample i ) {
    omega_i = random direction distributed according to pdf;
    L += shade( trace(x, omega_i) )
      * brdf * cos / pdf(omega_i);
  }
  return L /= N;
```

- When/how does the recursion stop*?
- Other problems?
 - You spent exponentially **more computation time** for effects that **contribute exponentially less** to the final image!

Exponential Explosion

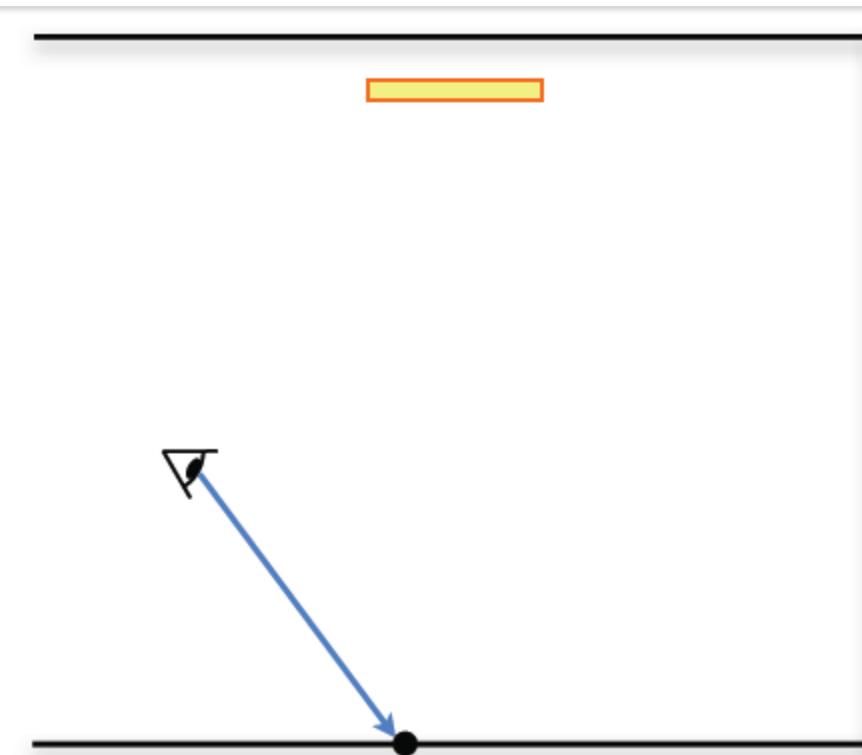
```
color shade( point x )
    if( x.Le != 0 ) return x.Le;
    color L = 0;
    for( each sample i ) {
        omega_i = random direction distributed according to pdf;
        L += shade( trace(x, omega_i) ) * brdf * cos / pdf(omega_i);
    }
    return L /= N;
```

- How can we avoid this exponential compute explosion?
 - Use an MC estimator with a single sample!

$$L(x, \omega_o) \approx L_e(x, \omega_o) + \frac{L(x', \omega) f_r(x, \omega, \omega_o) \cos \theta}{pdf(\omega)}$$

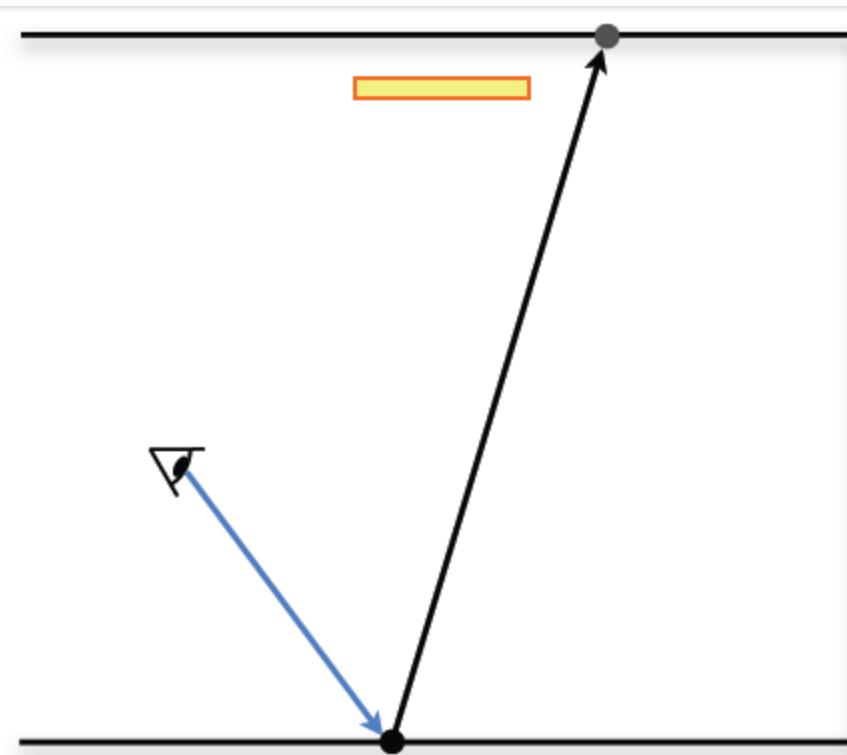
Path-tracing

$$L(x, \omega_o) \approx L_e(x, \omega_o) + \frac{L(x', \omega) f_r(x, \omega, \omega_o) \cos \theta}{pdf(\omega)}$$



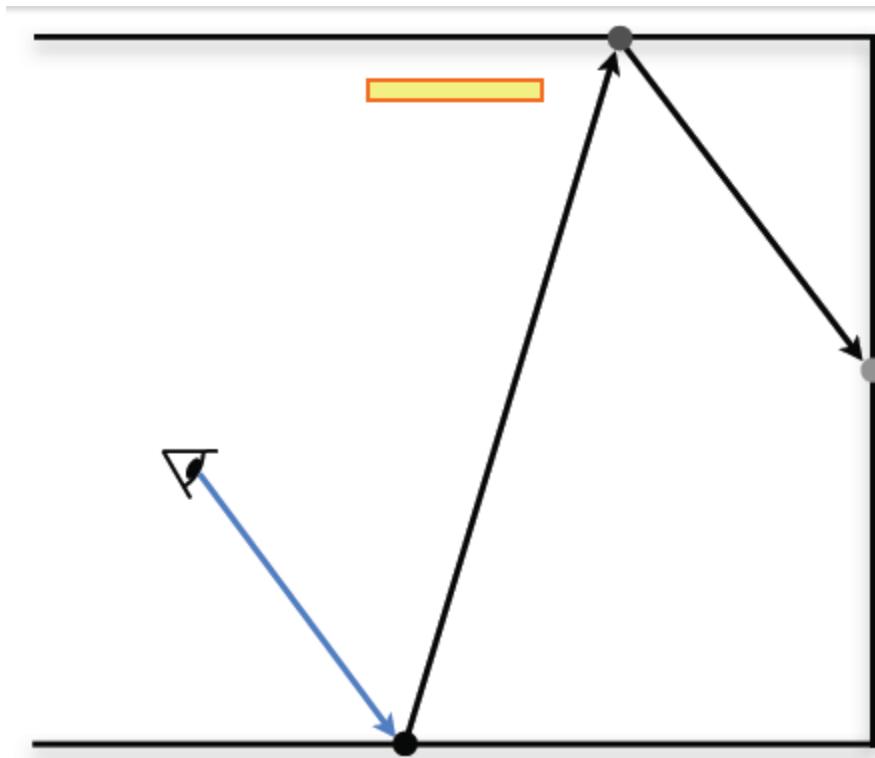
Path-tracing

$$L(x, \omega_o) \approx L_e(x, \omega_o) + \frac{L(x', \omega) f_r(x, \omega, \omega_o) \cos \theta}{pdf(\omega)}$$



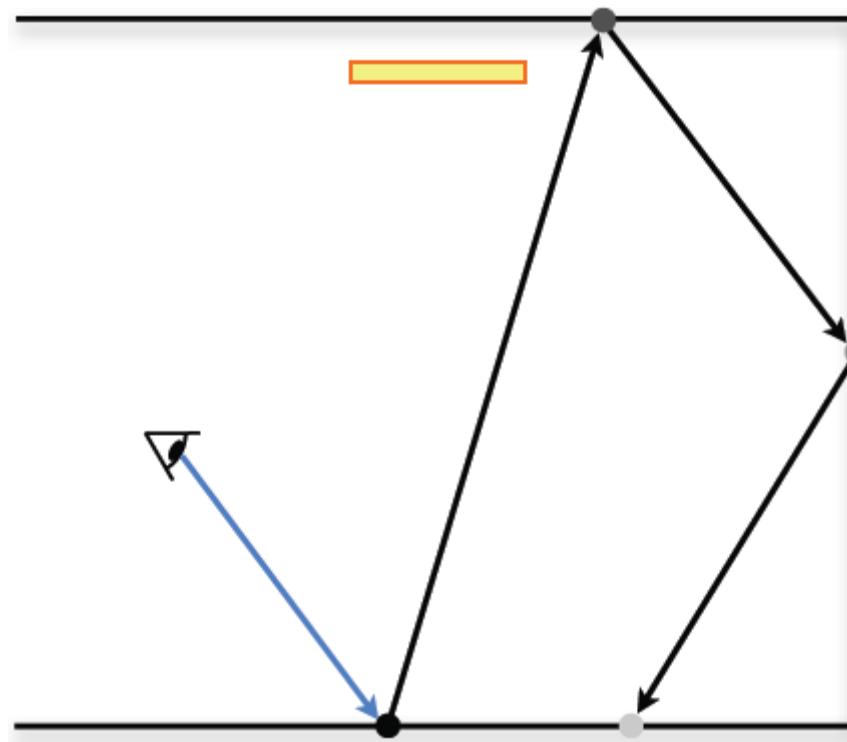
Path-tracing

$$L(x, \omega_o) \approx L_e(x, \omega_o) + \frac{L(x', \omega) f_r(x, \omega, \omega_o) \cos \theta}{pdf(\omega)}$$



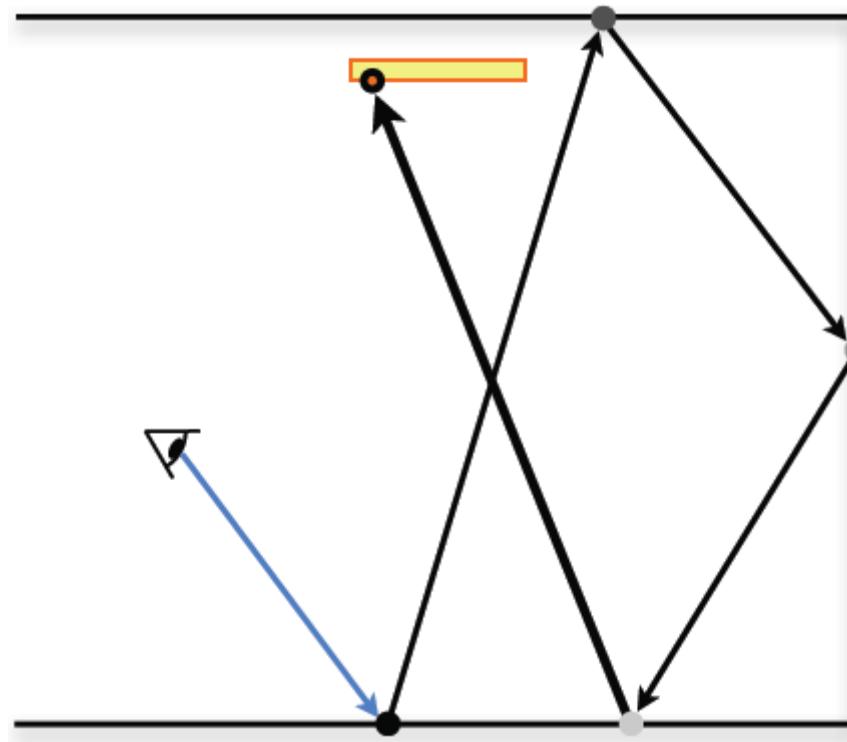
Path-tracing

$$L(x, \omega_o) \approx L_e(x, \omega_o) + \frac{L(x', \omega) f_r(x, \omega, \omega_o) \cos \theta}{pdf(\omega)}$$



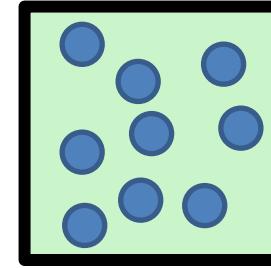
Path-tracing

$$L(x, \omega_o) \approx L_e(x, \omega_o) + \frac{L(x', \omega) f_r(x, \omega, \omega_o) \cos \theta}{pdf(\omega)}$$



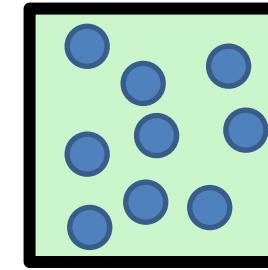
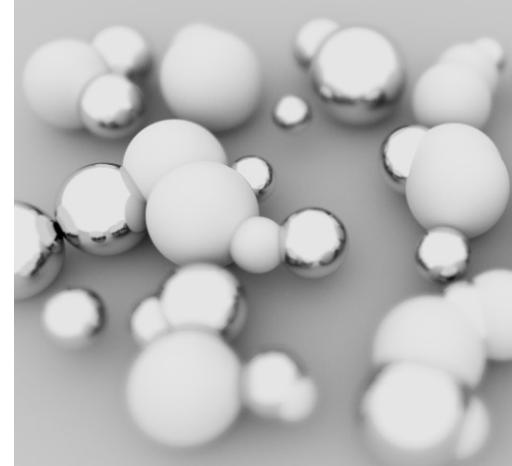
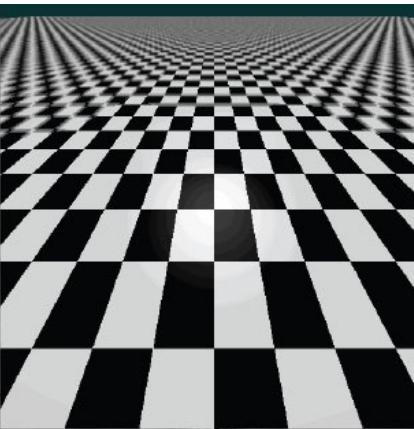
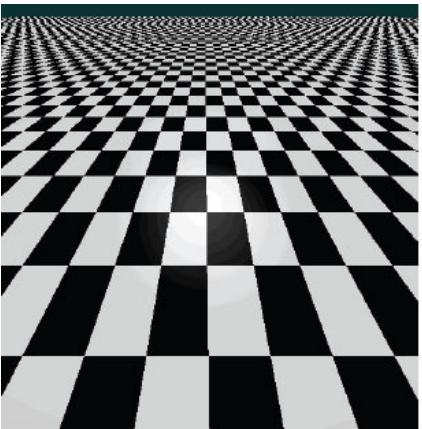
Implicit Path-tracing

```
color shade( point x )
    if( x.Le != 0 ) return x.Le;
    color L = 0;
    omega_i = random direction distributed according to pdf;
    L += shade( trace(x, omega_i) ) * brdf * cos / pdf(omega_i);
    return L;
```



- Then, in order to better approximate the integral, we just need to generate (and average) more *paths*
 - Shoot more rays through each pixel, each one creating a new path that contributes (potentially) to the pixel's final color
 - We can get anti-aliasing and other distribution effects “for free” when doing this pixel sampling

Distribution Effects



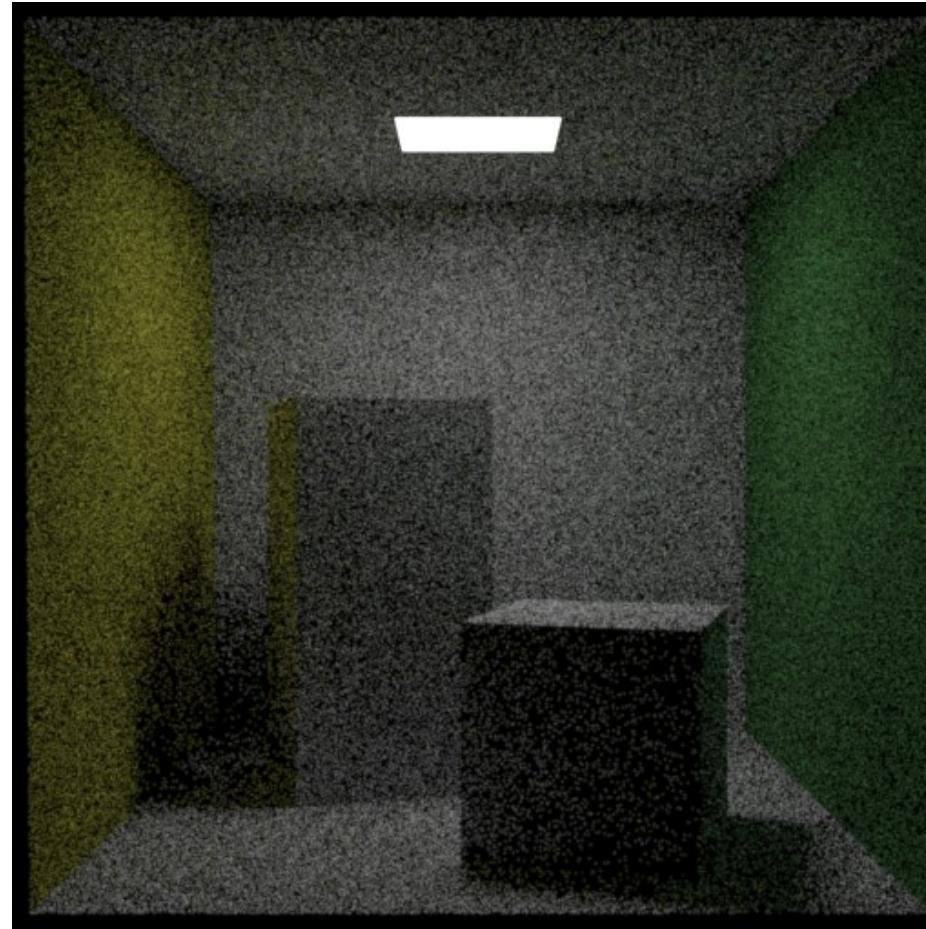
How do we stop?

Implicit Path-tracing: noisy results

$$L(x, \omega_o) \approx \frac{1}{N} \sum_N \left[L_e(x, \omega_o) + \frac{L(x', \omega) f_r(x, \omega, \omega_o) \cos \theta}{pdf(\omega)} \right]$$

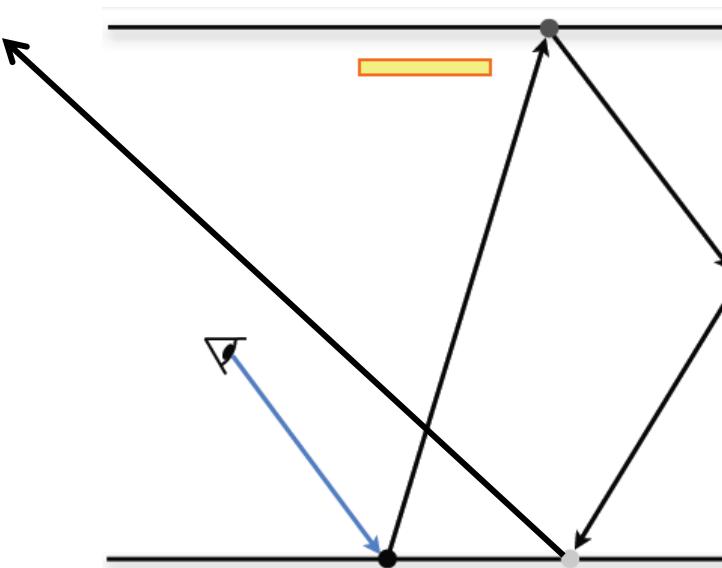
```
color shade( point x )
    if( x.Le != 0 ) return x.Le;
    color L = 0;
    omega_i = random direction
        sampled from pdf;
    L += shade( trace(x, omega_i) )
        * brdf * cos / pdf(omega_i);
    return L;

    ...
color c = 0;
for( each pixel sample s)
    c += shade(trace(eye, s))
c /= N;
    ...
```



Implicit Path-tracing: noisy results

- Why is there so much noise?
- The recursion terminates (with a non-zero value) only* if (the last vertex of) the path hits a light!
- Tons of wasted paths!



Split the Integral + Specialize

- By separately sampling the direct- and indirect-contributions of the rendering equation (with different strategies), we can reduce variance

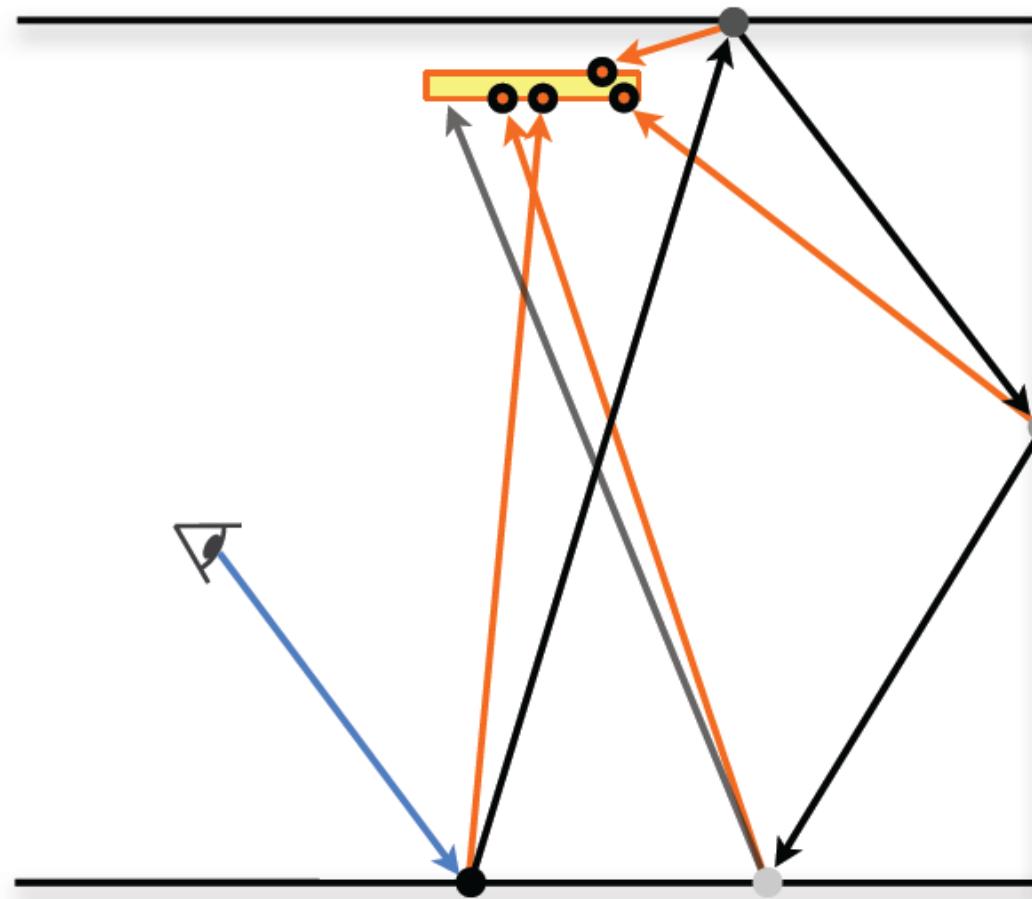
$$L = L_e(x, \omega_o) + L_d(x, \omega_o) + L_{ind}(x, \omega_o)$$

- Here, we'll still use a 1-sample estimator for each term
 - Direct lighting will be *explicitly* sampled
 - Need to be careful not to double count light...

$$\begin{aligned} &\approx L_e + \frac{L_e(x_d, \omega_d) f_r(x, \omega_i, \omega_d) V(x, \omega_d) \cos \theta_d \cos \theta_{x \rightarrow d}}{d_{x \rightarrow d}^2 \text{pdf}(x_d)} \\ &+ \frac{L(x', \omega) f_r(x, \omega, \omega_o) \cos \theta}{\text{pdf}(\omega)} \end{aligned}$$

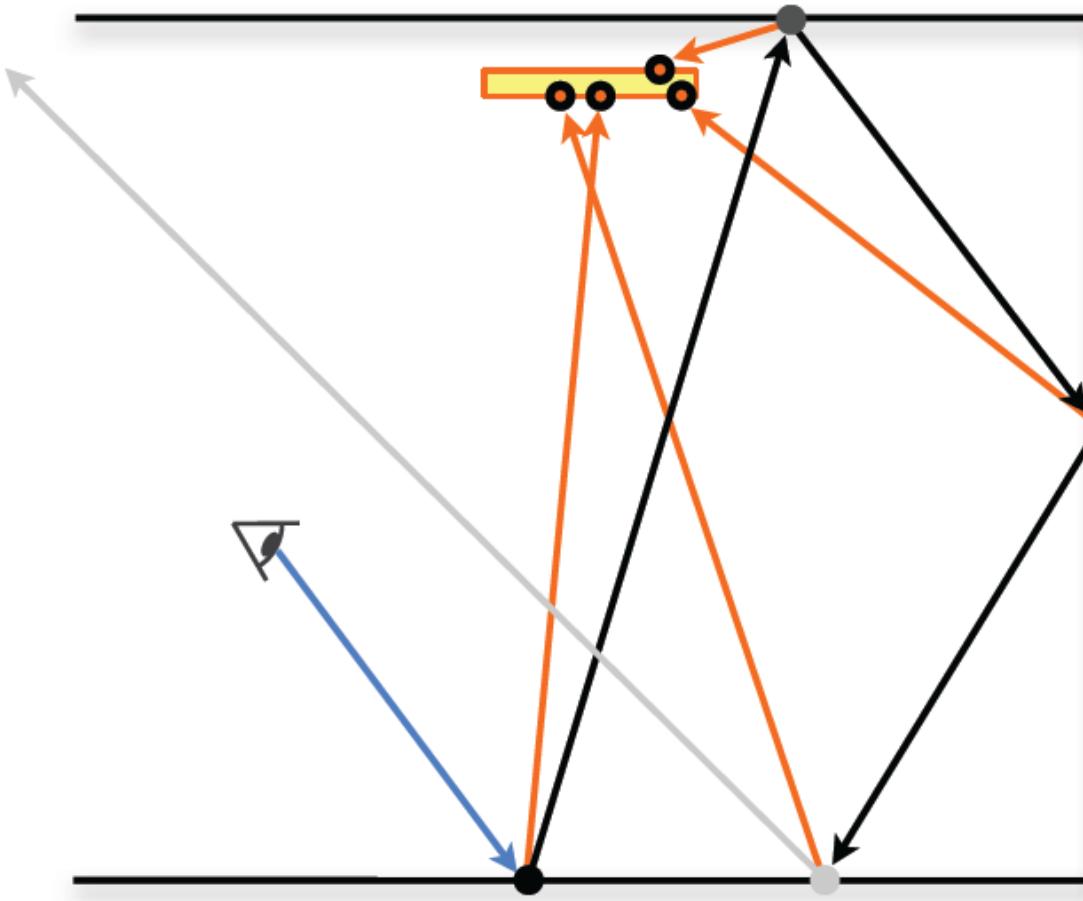
Explicit Path-tracing

(a.k.a. path-tracing with explicit direct lighting)

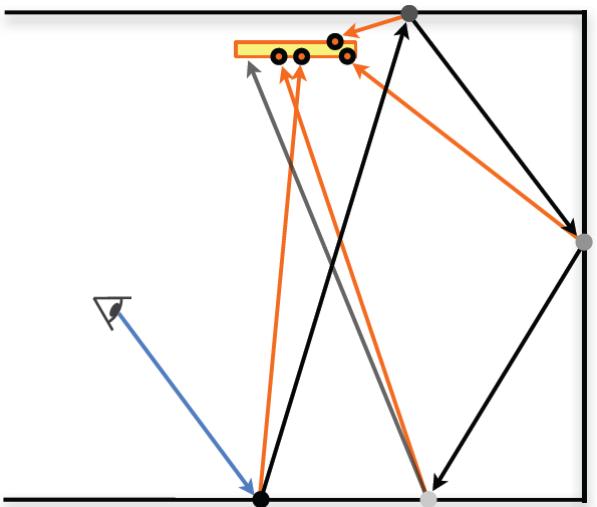


Explicit Path-tracing

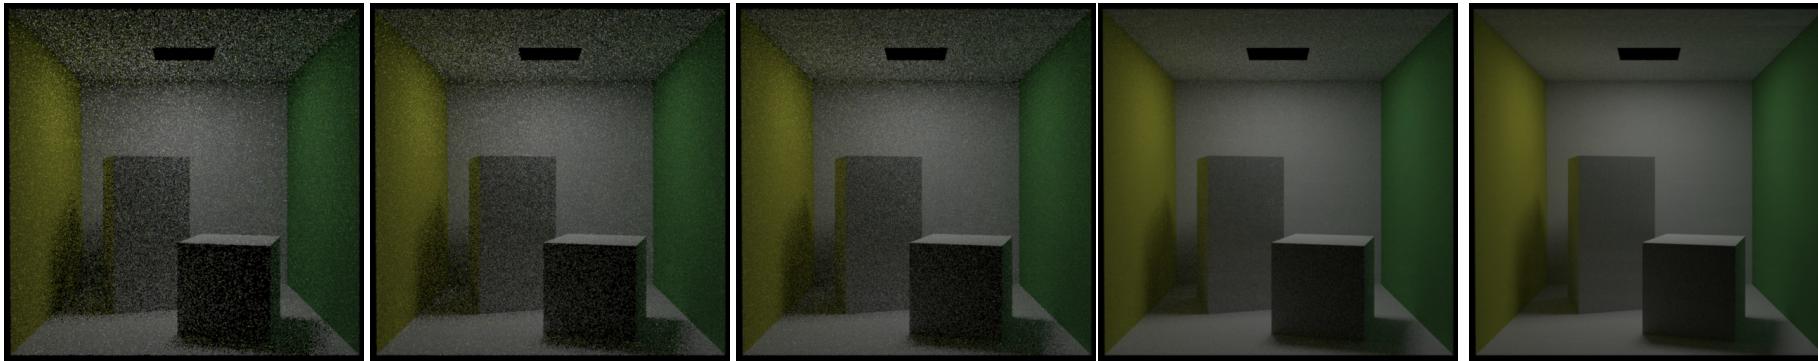
(a.k.a. path-tracing with explicit direct lighting)



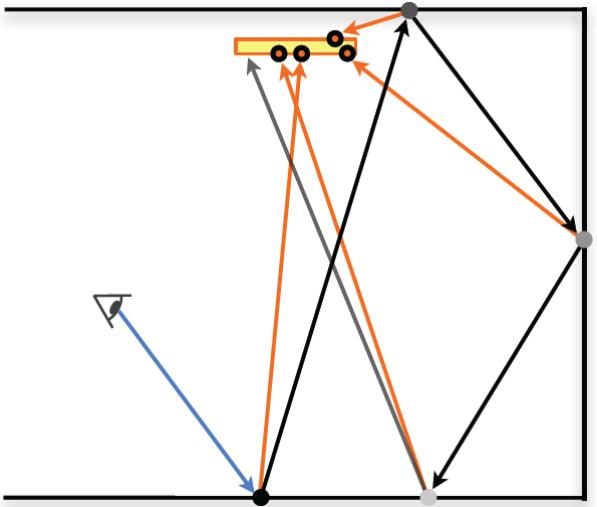
Explicit Path-tracing



```
color shade( point x )
color L = x.Le;
point p = random point on the light;
omega_i = random direction from pdf;
// L_d
if( trace(x, p) )
    L += fr * cos * p.Le * cos' / (|x-p|^2 * pdf_x)
// L_ind
L += shade(trace(x, omega_i)) * fr' * cos' / pdf;
return L;
```



Infinite Recursion

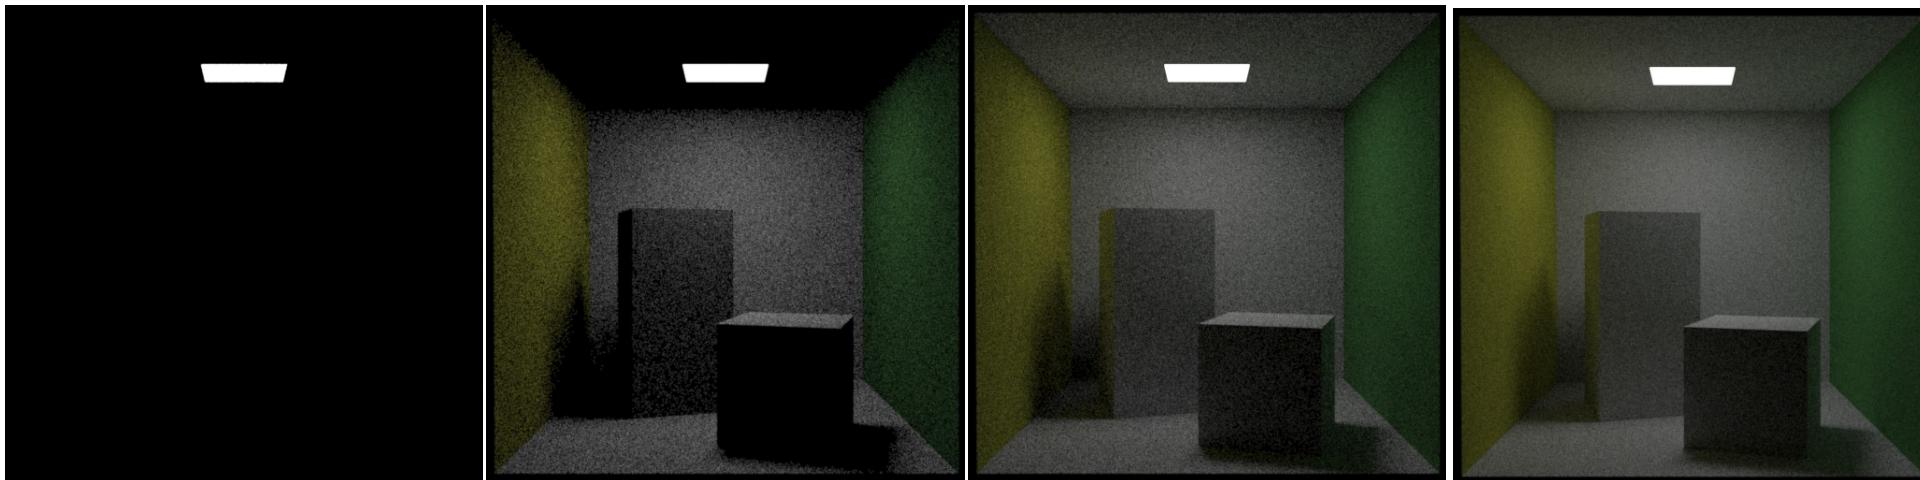


```
color shade( point x )  
color L = x.Le;  
point p = random point on the light;  
omega_i = random direction from pdf;  
// L_d  
if( trace(x, p) )  
    L += fr * cos * p.Le * cos' / (|x-p|^2 * pdf_x)  
// L_ind  
L += shade(trace(x, omega_i)) * fr' * cos' / pdf;  
return L;
```

There's still a (potentially) infinite recursion!

How to terminate the recursion?

- When/how do we stop the recursion?
 - Clamp with a maximum path length (recursion depth)?
 - This *biases* our result



p=1

p=2

p=3

p=4

- Are there unbiased ways to terminate recursion?

Russian Roulette

- Russian roulette is a probabilistic approach to terminate MC processes without changing their expected value / average (i.e., without biasing the result)
1. Pick some threshold / termination-probability χ
 2. Generate a random number ε between $(0,1]$
 3. Return $V' = \frac{V}{1-\chi}$ for a sample when $\varepsilon \leq \chi$ or 0 otherwise

$$E[V'] = (1 - \chi) \left[\frac{E[V]}{1 - \chi} \right] = E[V]$$

Russian Roulette

- To apply RR to path-tracing, how do we choose χ ?
 - Note, as with the *pdf*, no matter your choice the estimator will remain “correct”
 - But, as with the *pdf*, a “smarter” choice will improve convergence and reduce variance
 - e.g., for diffuse f_r , a typical choice is $\chi = \rho_x$
 - Interpretation: the higher the reflectance ρ_x , the higher the surface’s reflection, and so we would like to promote such reflections with higher probability

Explicit Path Tracing w/ Russian Roulette

- So, our final path tracing algorithm is:

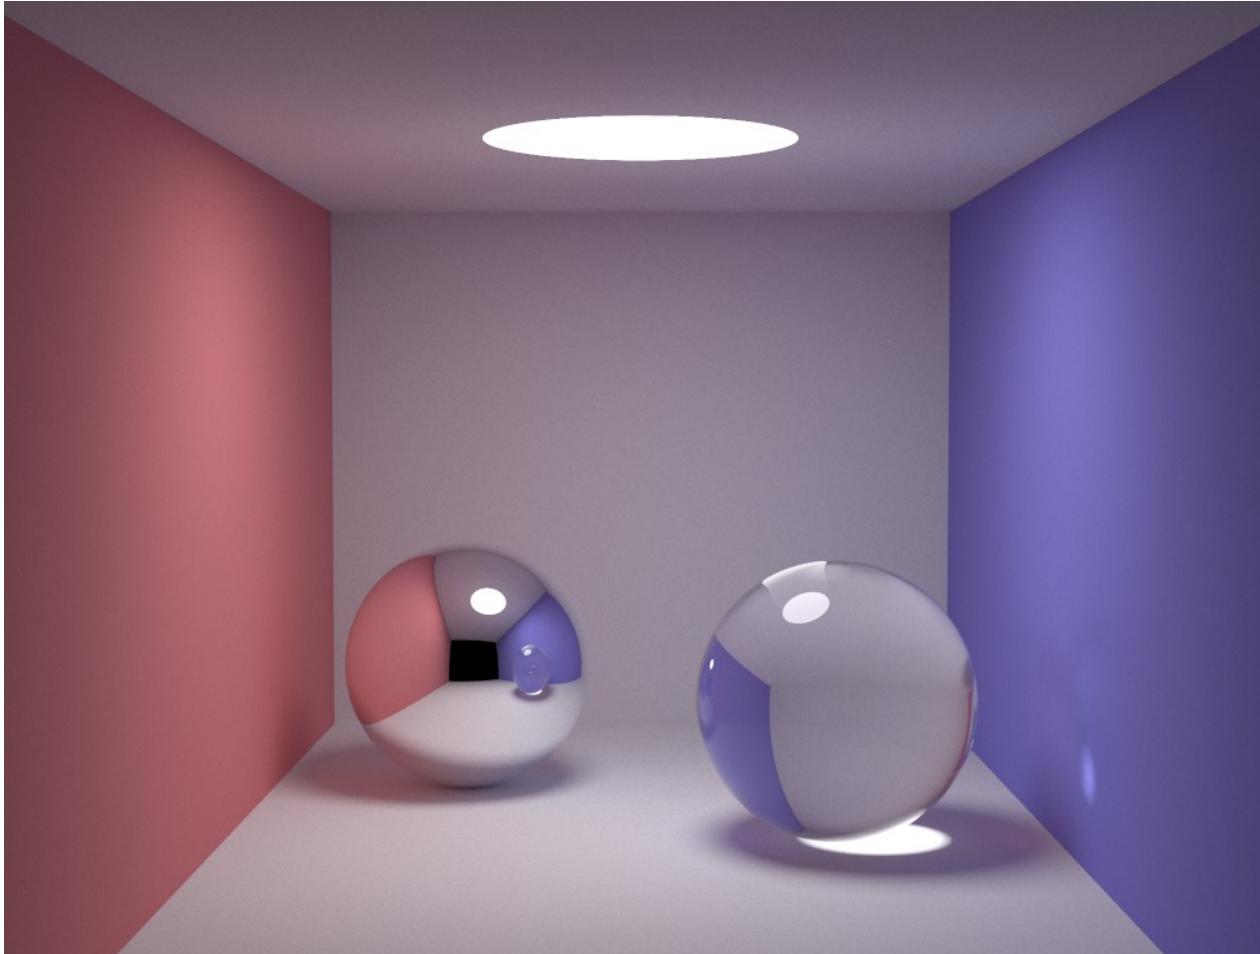
```
color shade( point x )
    // L_d
    point p = random point on the light;
    if( trace(x, p) )
        L += fr * cos * p.Le * cos' / (|x-p|^2 * pdf_x)

    // L_ind
    if( rand() < X )
        omega_i = random direction taken from pdf;
        L += shade(trace(x,omega_i)) * fr' * cos' / (pdf * (1-X));

    return L;
```

- With a basic ray tracer in place, we can code a basic *path tracer* in about 50 (very careful) lines of code!

Explicit Path Tracing w/ Russian Roulette



<http://www.kevinbeason.com/smallpt/>
99 lines of C++

Problems with Path Tracing

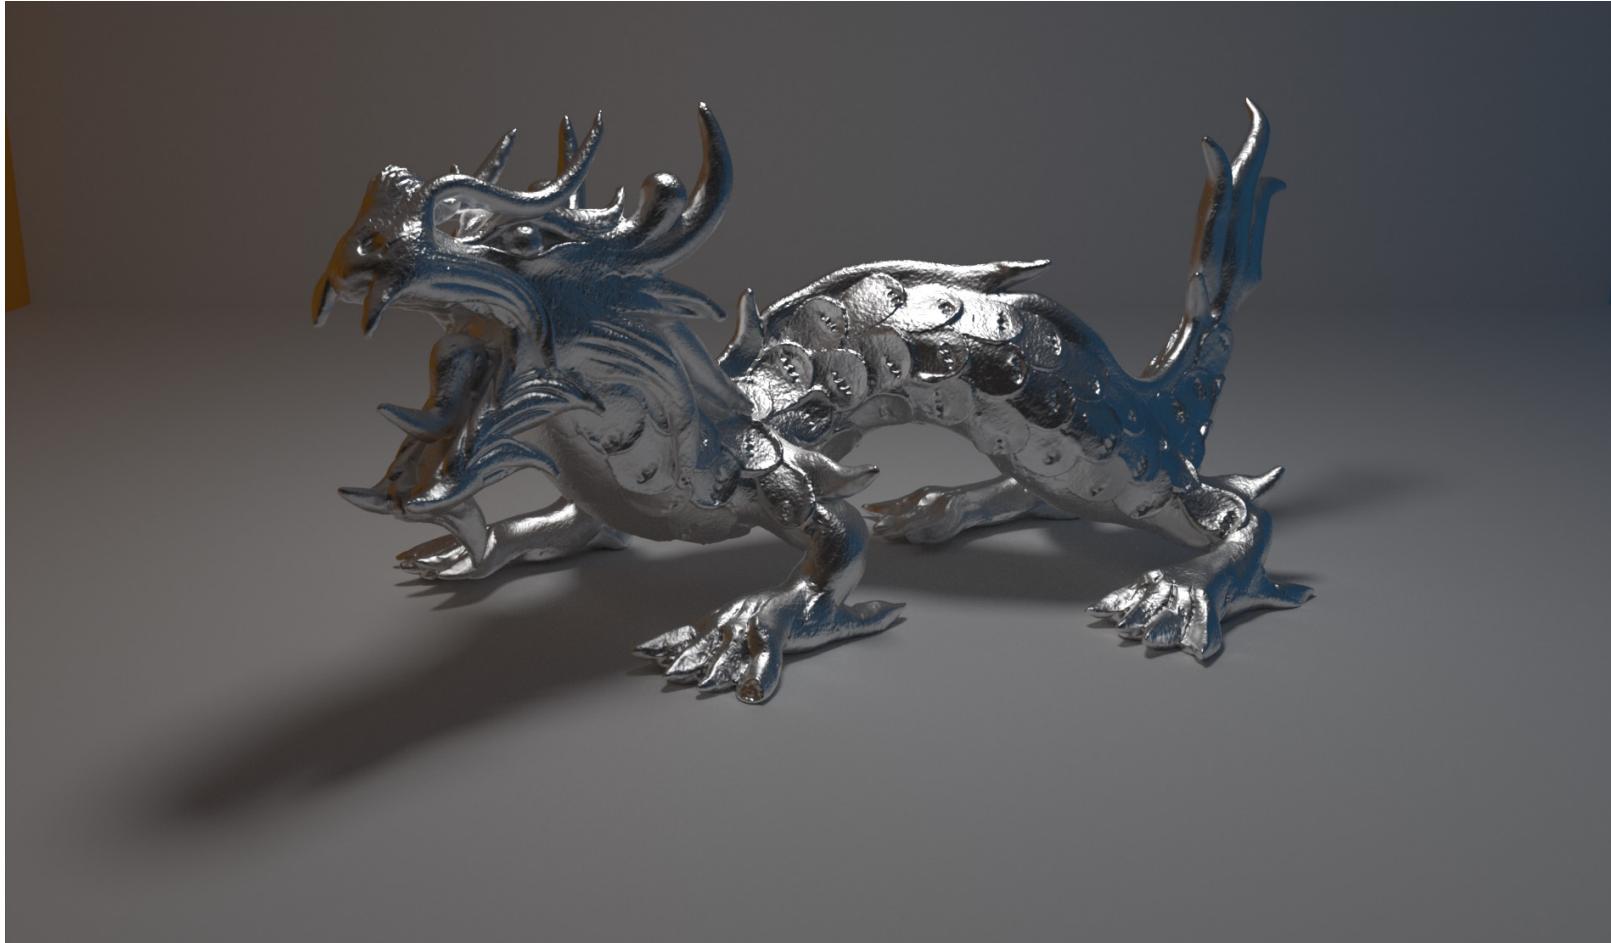
- Advantages:

- Easy to implement!
- Converges to the right answer!



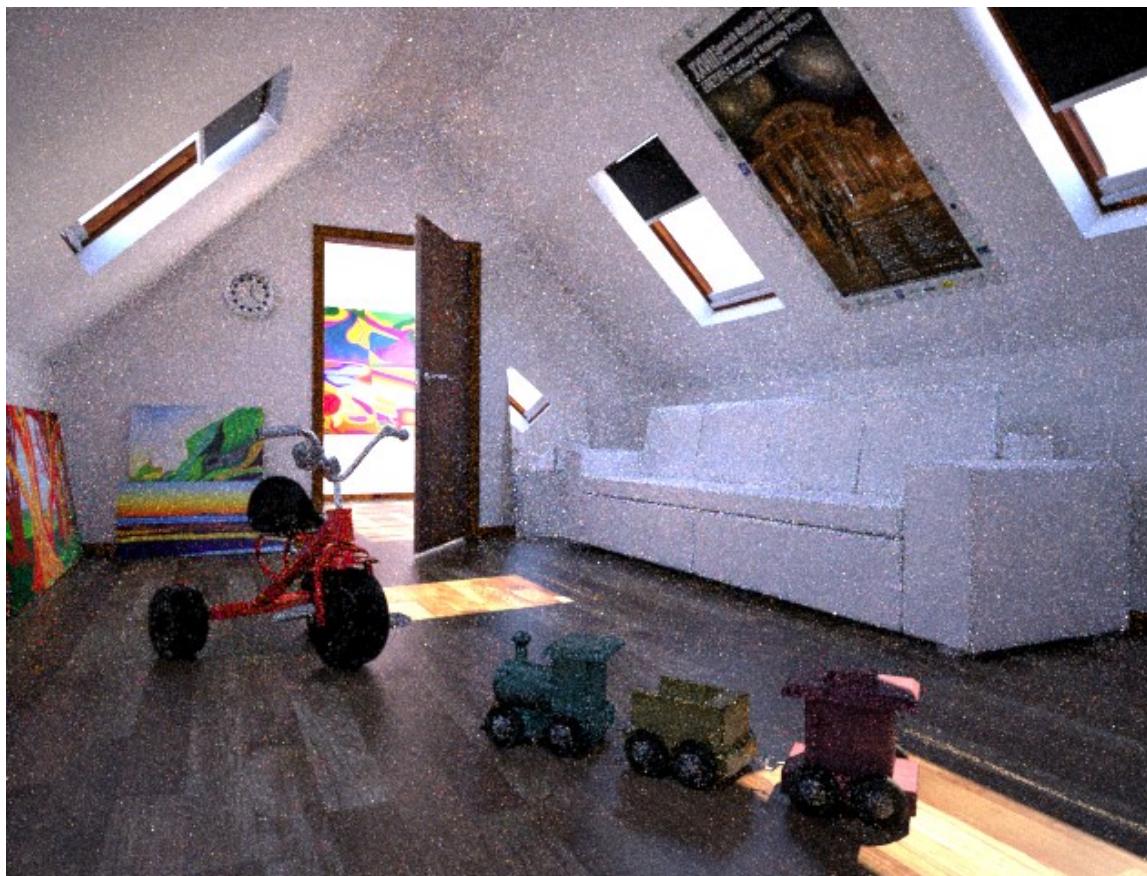
Problems with Path Tracing

- Converges to the right answer



Problems with Path Tracing

- Disadvantages:
 - S-L-O-W and noisy convergence



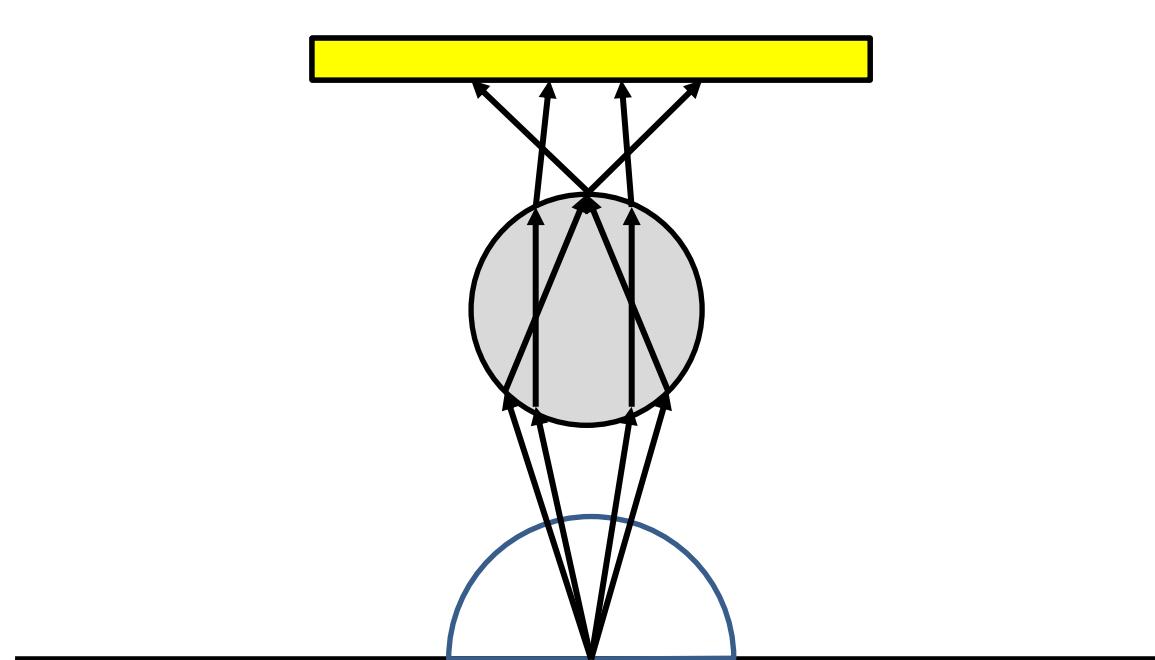
Problems with Path Tracing

– S-L-O-W and noisy convergence



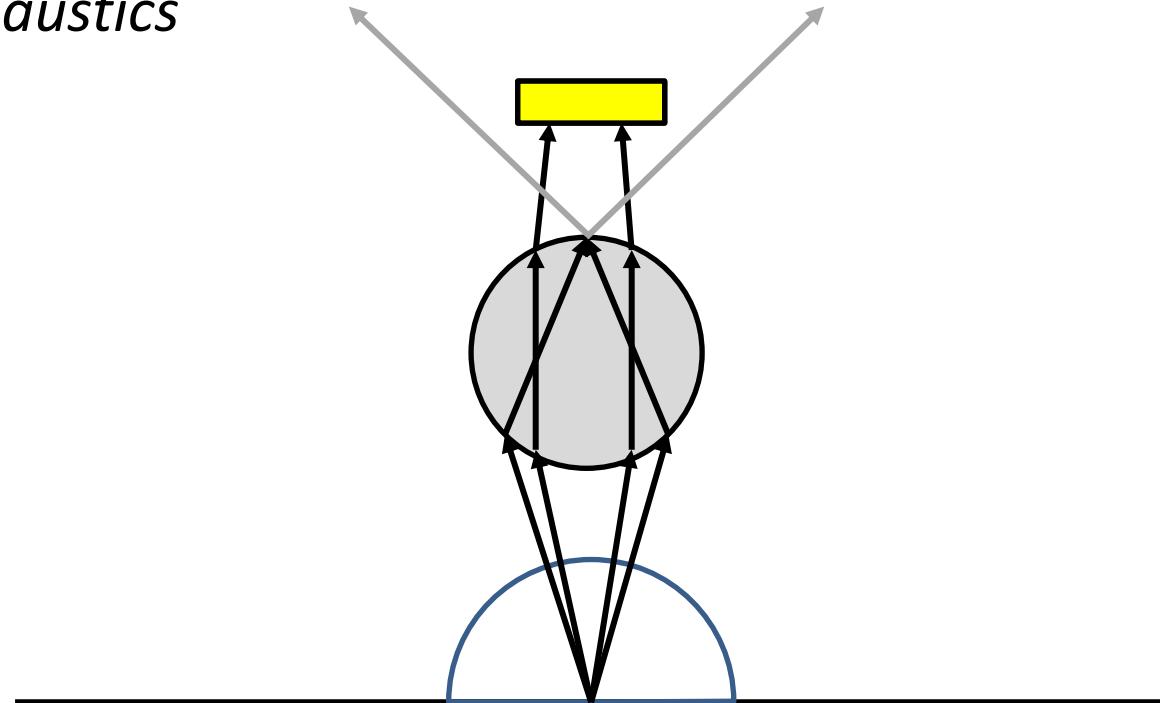
Problems with Path Tracing

- Disadvantages:
 - Difficulty handling complex illumination paths:
 - *Caustics (L S* D E)*



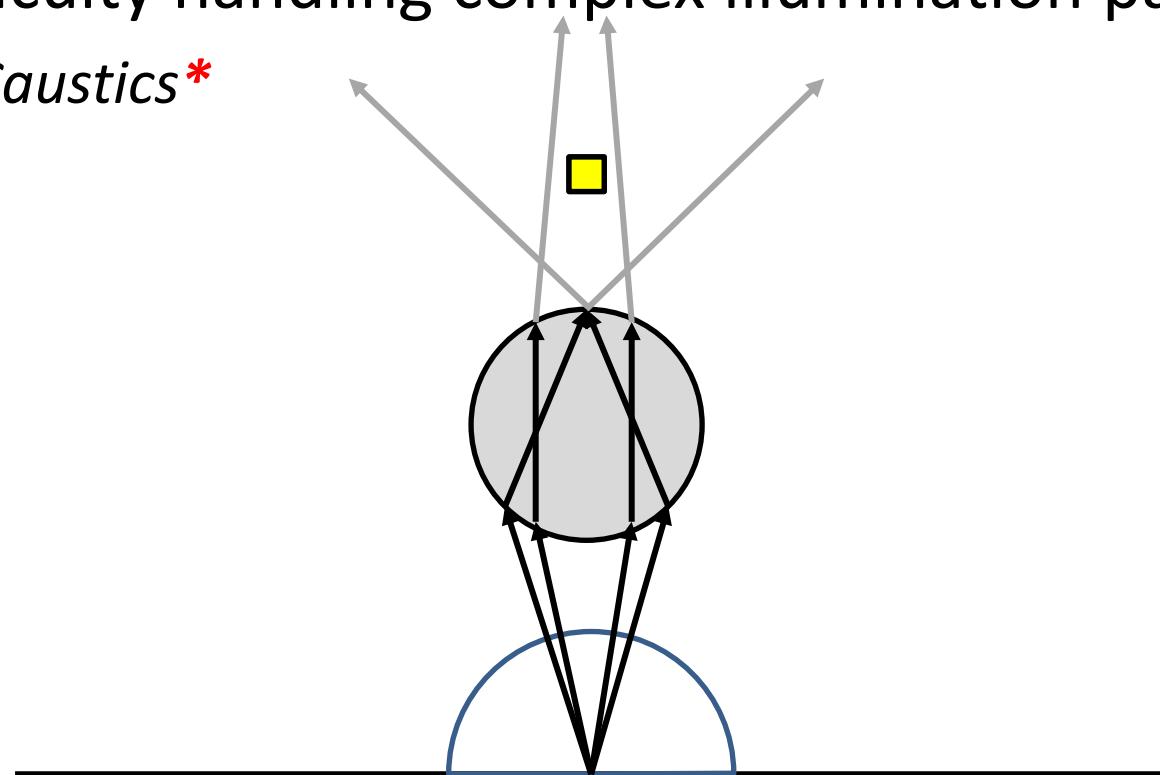
Problems with Path Tracing

- Disadvantages:
 - Difficulty handling complex illumination paths:
 - *Caustics*



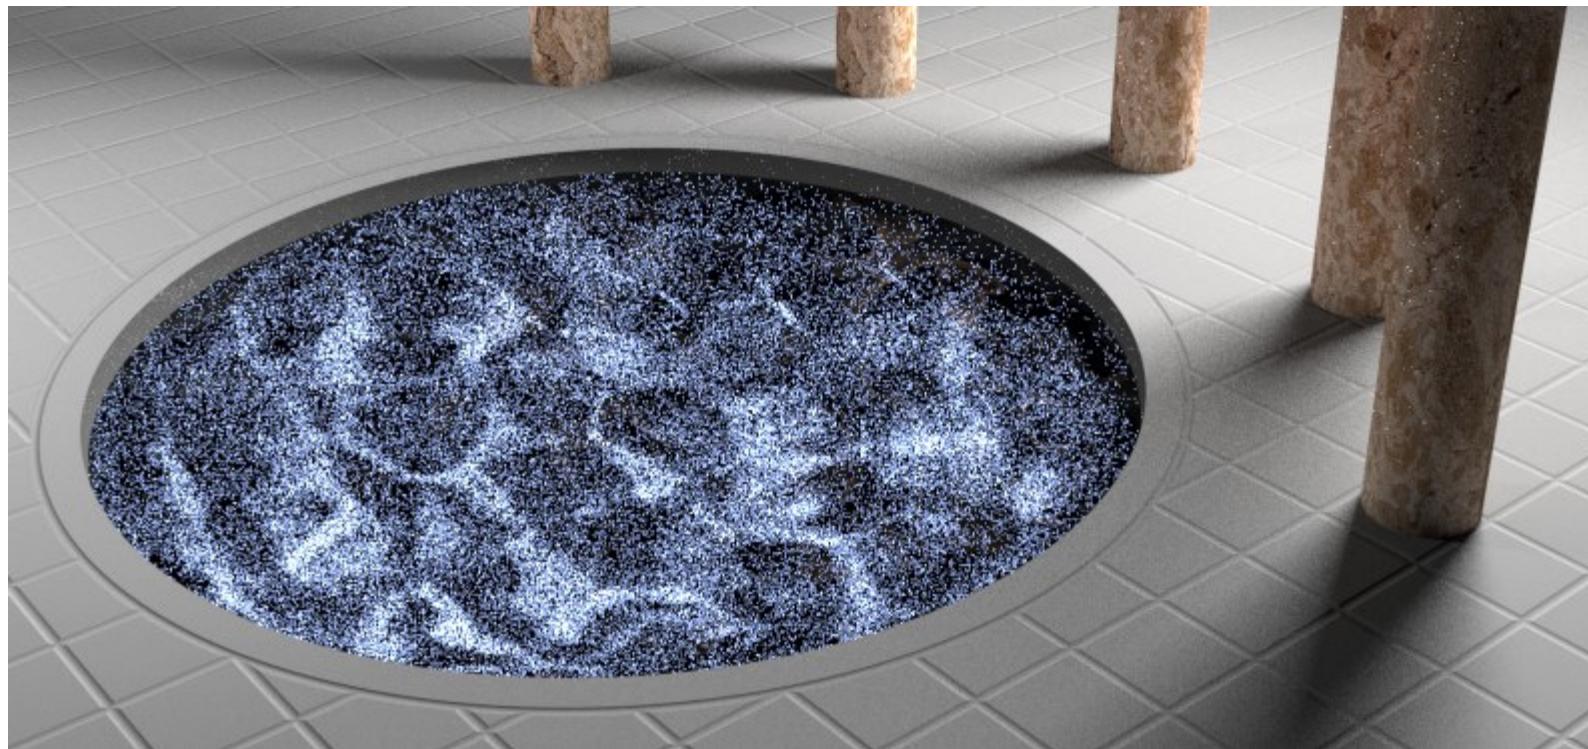
Problems with Path Tracing

- Disadvantages:
 - Difficulty handling complex illumination paths:
 - *Caustics**



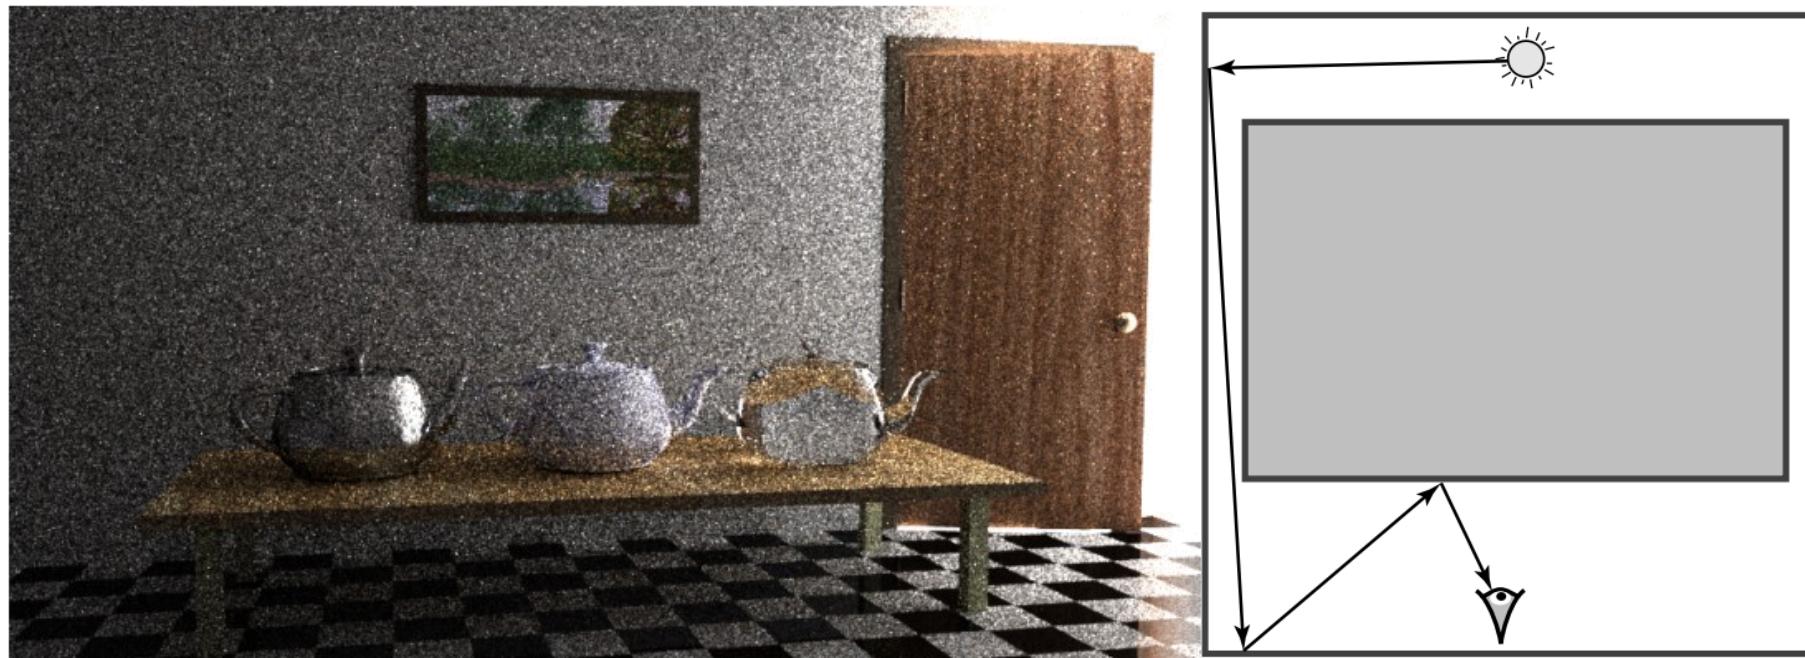
Problems with Path Tracing

- Disadvantages:
 - Difficulty handling complex illumination paths:
 - *Caustics*



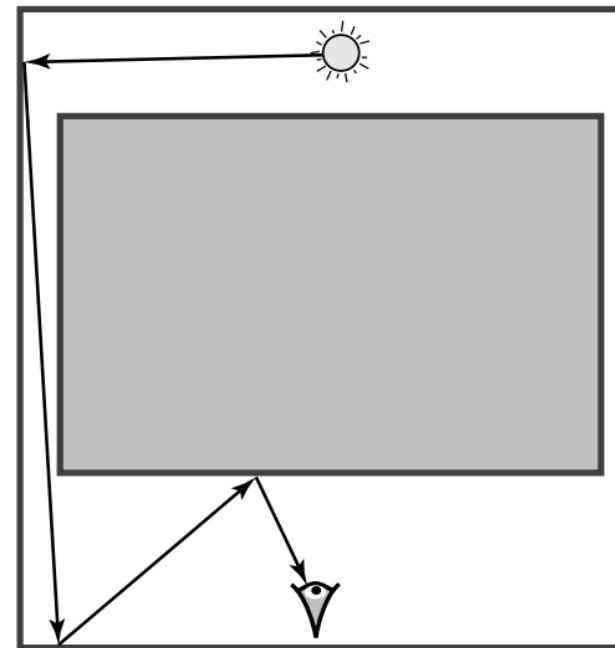
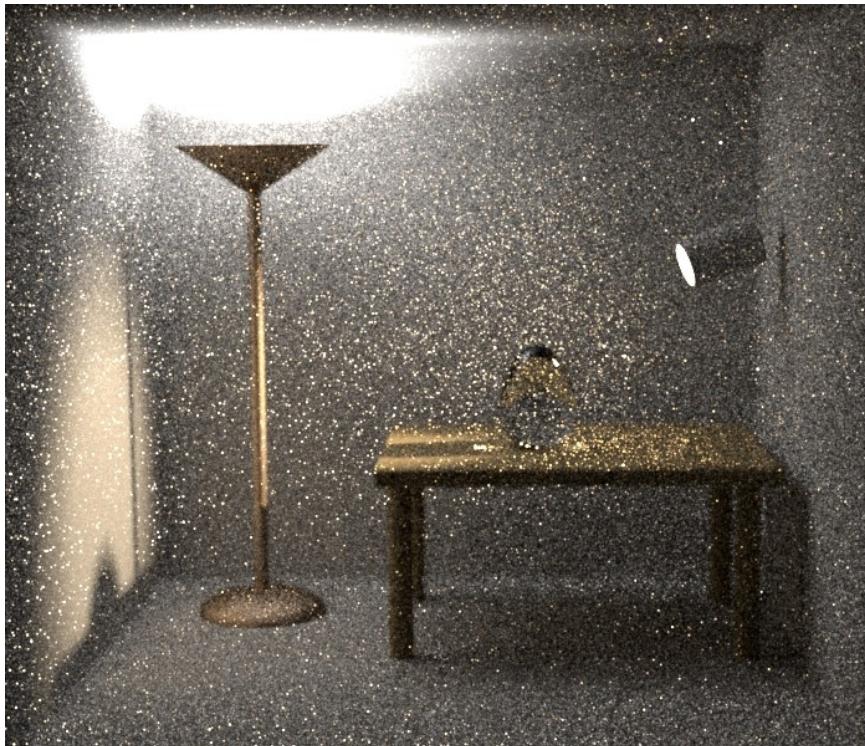
Problems with Path Tracing

- Disadvantages:
 - Difficulty handling complex illumination paths:
 - Lighting contributions primarily due to indirect light

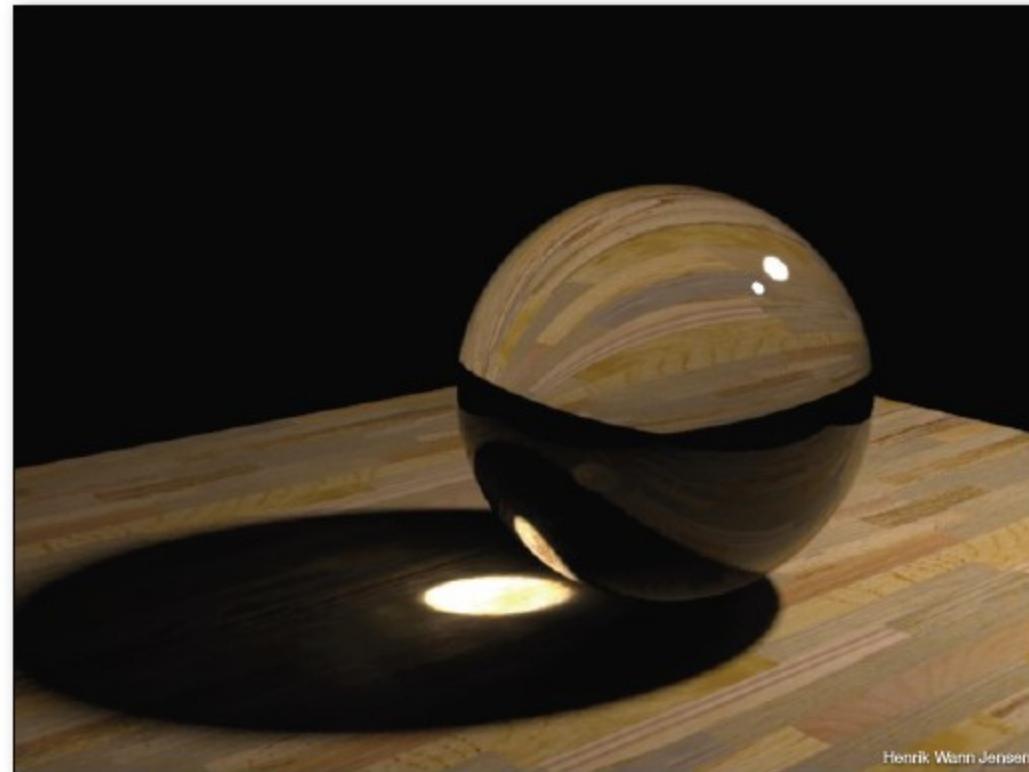


Problems with Path Tracing

- Disadvantages:
 - Difficulty handling complex illumination paths:
 - Lighting contributions primarily due to indirect light



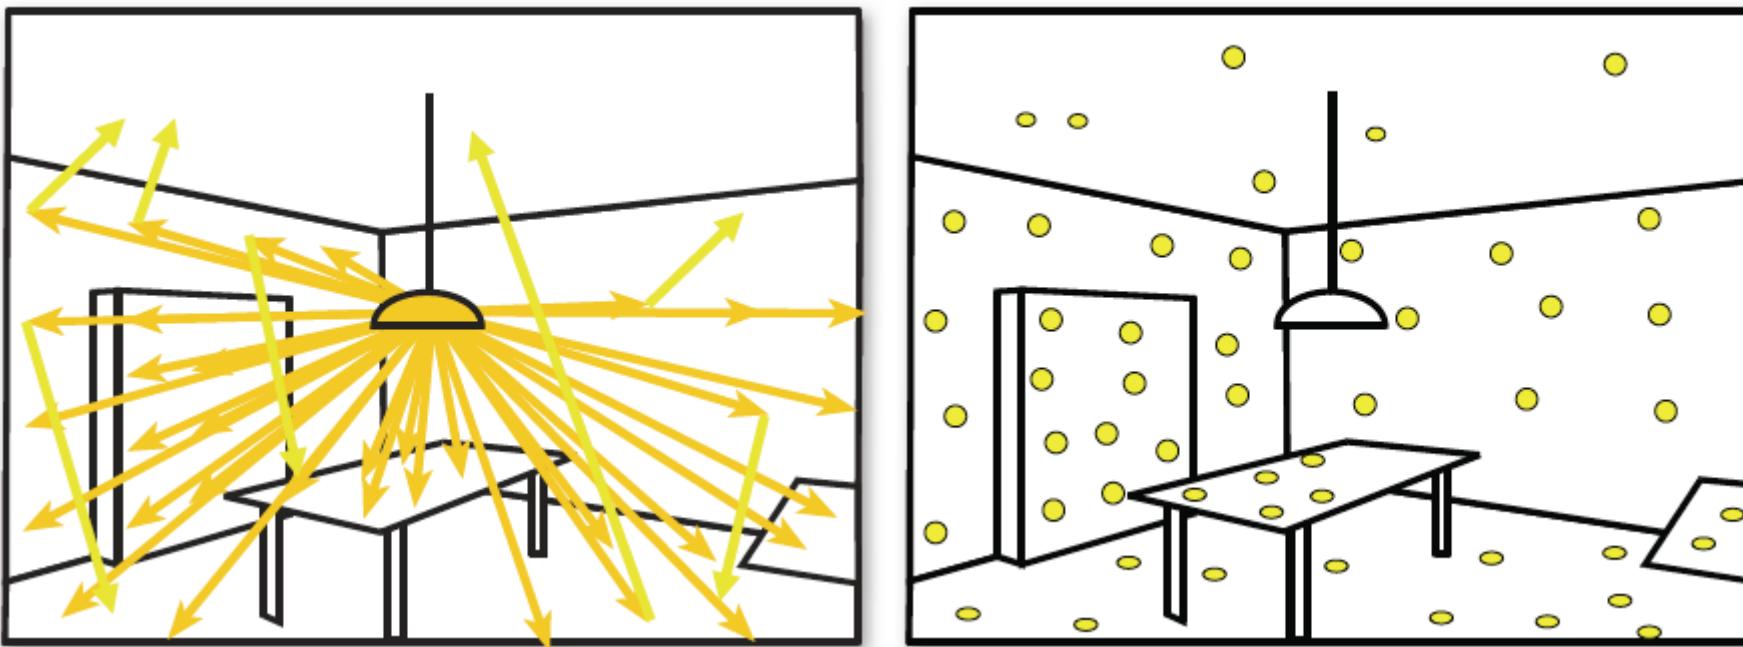
Photon Mapping



Henrik Wann Jensen

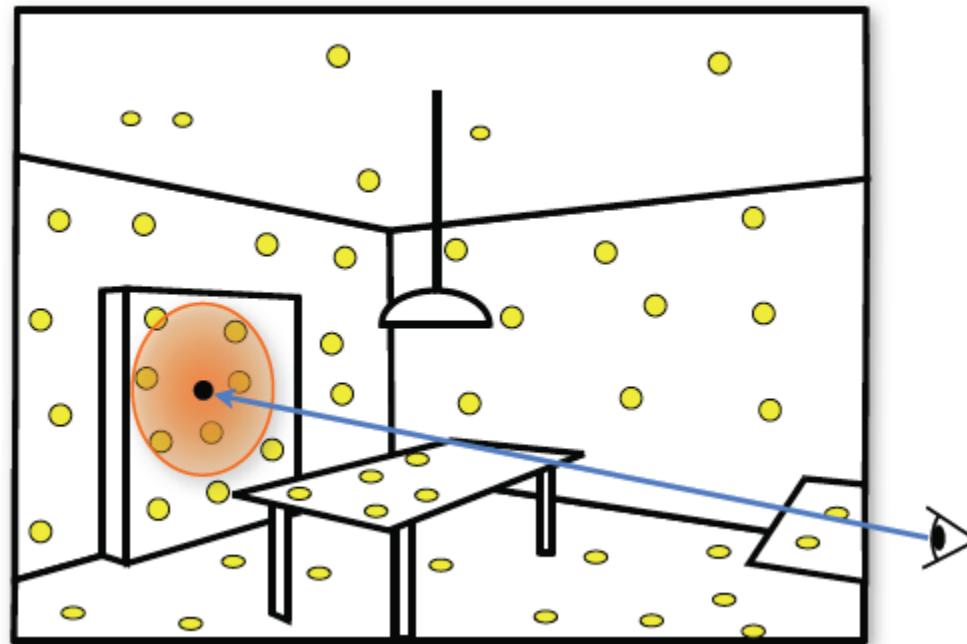
Photon Mapping

- Basic Photon Mapping is a two-step algorithm:
 1. Trace, reflect and deposit light particles (*photons*)



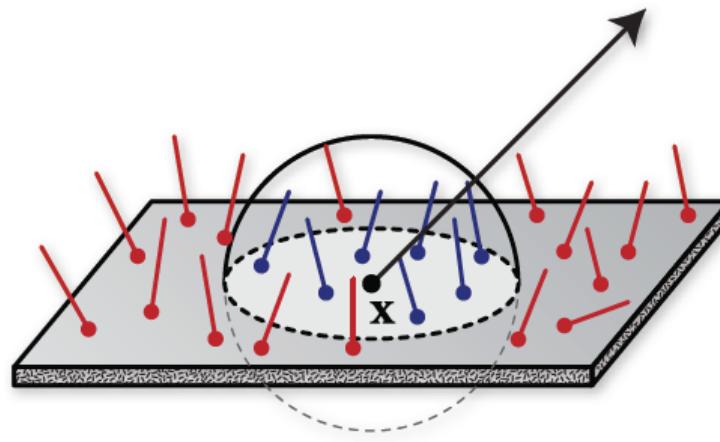
Photon Mapping

- Basic Photon Mapping is a two-step algorithm:
 1. Trace, reflect and deposit light particles (*photons*)
 2. Estimate light particle density to render image



Photon Mapping

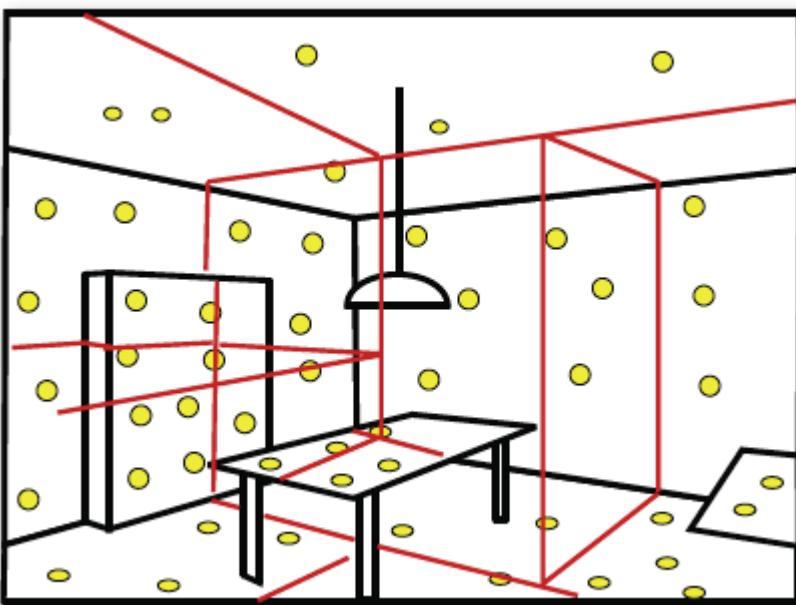
2. Estimate light particle density to render image



$$L(x, \omega_o) \approx \sum_{p=1}^N \frac{\Phi_p(x, \omega_p)}{\Delta A} f_r(x, \omega_p, \omega_o) \cos \theta_p$$

Photon Mapping

- We typically store the photons in a *kd-tree* in order to do quick lookups necessary for efficient density estimation



Henrik Wann Jensen

Photon Mapping



500000 photons / 500 photons in radiance estimate

Review and More Information

- Fundamentals of Computer Graphics
 - Section 24.4 Path Tracing (brief)
- Physically Based Rendering, from theory to implementation
 - <http://www.pbr-book.org/3ed-2018/contents.html>
 - <https://www.pbrt.org/>
 - Chapter 13 covers Monte Carlo Integration
 - Chapter 14 includes details on Path tracking (section 14.5)