




Images and Gamma

Digital camera

- Color typically captured using color mosaic

Mosaic Capture



[Foveon]

Raster image representation

- All these devices suggest 2D arrays of numbers
- Big advantage: represent arbitrary images
 - approximate arbitrary functions with increasing resolution
 - works because memory is cheap (brute force approach!)



[Philip Greenspun]

Meaning of a raster image

- Meaning of a given array is a function on 2D
- Define meaning of array = result of output device?
 - that is, piecewise constant for LCD, blurry for CRT
 - but: we don't have just one output device
 - but: want to define images we can't display (e.g. too big)
- Abstracting from device, problem is reconstruction
 - image is a sampled representation
 - pixel means "this is the intensity around here"
 - LCD: intensity is constant over square regions
 - CRT: intensity varies smoothly across pixel grid
 - will discuss specifics of reconstruction later

Datatypes for raster images

- Bitmaps: `boolean` per pixel (1 bpp): $I : \mathbb{R}^2 \rightarrow \{0, 1\}$
 - interp. = black and white; e.g. fax
- Grayscale: integer per pixel: $I : \mathbb{R}^2 \rightarrow [0, 1]$
 - interp. = shades of gray; e.g. black-and-white print
 - precision: usually `byte` (8 bpp); sometimes 10, 12, or 16 bpp
- Color: 3 integers per pixel: $I : \mathbb{R}^2 \rightarrow [0, 1]^3$
 - interp. = full range of displayable color; e.g. color print
 - precision: usually `byte[3]` (24 bpp)
 - sometimes 16 (5+6+5) or 30 or 36 or 48 bpp
 - indexed color: a fading idea

Datatypes for raster images

- Floating point: $I : \mathbb{R}^2 \rightarrow \mathbb{R}_+$ or $I : \mathbb{R}^2 \rightarrow \mathbb{R}_+^3$
 - more abstract, because no output device has infinite range
 - provides *high dynamic range* (HDR)
 - represent real scenes independent of display
 - becoming the standard intermediate format in graphics processors
- Clipping and white point
 - common to compute FP, then convert to integer
 - full range of values may not “fit” in display’s output range
 - simplest solution: choose a maximum value, scale so that value becomes full intensity ($2^n - 1$ in an n -bit integer image)



exposure:
-8 stops

image: Paul Debevec

exposure:
+0 stops



image: Paul Debevec


exposure:
+6 stops



image: Paul Debevec

Datatypes for raster images

- For color or grayscale, sometimes add *alpha* channel
 - describes transparency of images
 - more on this in a few lectures



Storage requirements for images

- 1024x1024 image (1 megapixel)
 - bitmap: 128KB
 - grayscale 8bpp: 1MB
 - grayscale 16bpp: 2MB
 - color 24bpp: 3MB
 - floating-point HDR color: 12MB



Converting pixel formats

- Color to gray
 - could take one channel (blue, say)
 - leads to odd choices of gray value
 - combination of channels is better
 - but different colors contribute differently to lightness
 - which is lighter, full blue or full green?
 - good choice: $\text{gray} = 0.2 R + 0.7 G + 0.1 B$
 - more on this in color, later on

Same pixel values.



Same luminance?



Converting pixel precision

- Up is easy; down loses information—be careful



[photo: Philip Greenspun]

1 bpp (2 grays)

Dithering

- When decreasing bpp, we quantize
- Make choices consistently: banding
- Instead, be inconsistent—dither
 - turn on some pixels but not others in gray regions
 - a way of trading spatial for tonal resolution
 - choose pattern based on output device
 - laser, offset: clumped dots required (halftone)
 - inkjet, screen: dispersed dots can be used

Dithering methods

- Ordered dither
 - based on traditional, optically produced halftones
 - produces larger dots
- Diffusion dither
 - takes advantage of devices that can reproduce isolated dots
 - the modern winner for desktop printing



Gamma and Displays

Intensity encoding in images


- What do the numbers in images (pixel values) mean?
 - they determine how bright that pixel is
 - bigger numbers are (usually) brighter
- *Transfer function*: function that maps input pixel value to luminance of displayed image

$$I = f(n) \quad f : [0, N] \rightarrow [I_{\min}, I_{\max}]$$

- What determines this function?
 - physical constraints of device or medium
 - desired visual characteristics

Transfer Function

- Something like this:



Constraints on transfer function


- Maximum displayable intensity, I_{\max}
 - how much power can be channeled into a pixel?
 - LCD: backlight intensity, transmission efficiency (<10%)
 - projector: lamp power, efficiency of imager and optics
- Minimum displayable intensity, I_{\min}
 - light emitted by the display in its “off” state
 - *e.g.* stray electron flux in CRT, polarizer quality in LCD
- Viewing flare, k : light reflected by the display
 - very important factor determining image contrast in practice
 - 5% of I_{\max} is typical in a normal office environment [sRGB spec]
 - much effort to make very black CRT and LCD screens
 - all-black decor in movie theaters

Dynamic range

- Dynamic range $R_d = I_{\max} / I_{\min}$, or $(I_{\max} + k) / (I_{\min} + k)$
 - determines the degree of image contrast that can be achieved
 - a major factor in image quality
- Ballpark values
 - Desktop display in typical conditions: 20:1
 - Photographic print: 30:1
 - Desktop display in good conditions: 100:1
 - Photographic transparency (directly viewed): 1000:1
 - High dynamic range display: 10,000:1

Transfer function shape

- Desirable property: the change from one pixel value to the next highest pixel value should not produce a visible contrast
 - otherwise smooth areas of images will show visible bands
- What contrasts are visible?
 - rule of thumb: under good conditions we can notice a 2% change in intensity
 - therefore we generally need smaller quantization steps in the darker tones than in the lighter tones
 - most efficient quantization is logarithmic



[Philip Greenspun]

an image with severe *banding*

How many levels are needed?

- Depends on dynamic range
 - 2% steps are $0 \mapsto I_{\min}; 1 \mapsto 1.02I_{\min}; 2 \mapsto (1.02)^2 I_{\min}; \dots$
 - $\log 1.02$ is about $1/120$, so 120 steps per decade of dynamic range
 - 240 for desktop display
 - 360 to print to film
 - 480 to drive HDR display
- If we want to use linear quantization (equal steps)
 - one step must be $< 2\%$ ($1/50$) of I_{\min}
 - need to get from ~ 0 to $I_{\min} \bullet R_d$ so need about $50 R_d$ levels
 - 1500 for a print; 5000 for desktop display; 500,000 for HDR display
- Moral: 8 bits is just barely enough for low-end applications
 - but only if we are careful about quantization


Intensity quantization in practice

- Option 1: linear quantization $I(n) = (n/N) I_{\max}$
 - pro: simple, convenient, amenable to arithmetic
 - con: requires more steps (wastes memory)
 - need 12 bits for any useful purpose; more than 16 for HDR
- Option 2: power-law quantization $I(n) = (n/N)^\gamma I_{\max}$
 - pro: fairly simple, approximates ideal exponential quantization
 - con: need to linearize before doing pixel arithmetic
 - con: need to agree on exponent
 - 8 bits are OK for many applications; 12 for more critical ones
- Option 2: floating-point quantization $I(x) = (x/w) I_{\max}$
 - pro: close to exponential; no parameters; amenable to arithmetic
 - con: definitely takes more than 8 bits
 - 16-bit “half precision” format is becoming popular

Why gamma?

- Power-law quantization, or *gamma correction* is most popular
- Original reason: CRTs are like that
 - intensity on screen is proportional to (roughly) voltage²
- Continuing reason: inertia + memory savings
 - inertia: gamma correction is close enough to logarithmic that there's no sense in changing
 - memory: gamma correction makes 8 bits per pixel an acceptable option

Gamma quantization



- Close enough to ideal perceptually uniform exponential

Gamma correction

- Sometimes (often, in graphics) we have computed intensities a that we want to display linearly
- In the case of an ideal monitor with zero black level,


$$I(n) = (n/N)^\gamma$$

(where $N = 2^n - 1$ in n bits). Solving for n :


$$n = N a^{\frac{1}{\gamma}}$$

- This is the “gamma correction” recipe that has to be applied when computed values are converted to 8 bits for output
 - failing to do this (implicitly assuming gamma = 1) results in dark, oversaturated images


Gamma correction



corrected for
©lower than
display



OK



corrected for
©higher than
display

[Philip Greenspun]

Review and More Information

- Fundamentals of Computer Graphics
 - Section 3.2.2 Monitor Intensities and Gamma